

Master Thesis

**A Haskell Web Application for Data Mining Competition Results
from StarExec**

A thesis submitted to attain the degree of
Master of Science at the Leipzig University of Applied Sciences
(M.Sc. HTWK Leipzig)

presented by
Stefan von der Krone

B.Eng. in Media Technology, University of applied Sciences Mittweida

born on 19th of April, 1984

citizen of Germany

thesis advisor: Prof. Dr. Johannes Waldmann

January 23rd, 2015

Table of Contents

Abstract	i
Preface	iii
1 Introduction	1
2 Termination	5
2.1 Termination of Term Rewriting	5
2.2 Termination Community	7
2.3 Termination Competition	8
2.4 TPDB - Termination Problem Data Base	9
3 StarExec	11
3.1 History	11
3.2 Data Model	12
3.3 Interfaces	13
4 Use Cases	17
4.1 Start a predefined Competition	17
4.2 Displaying Results and Outputs	18
4.3 Compare Results, Querying	20

5	Requirements	21
5.1	Solver and Post processor output	21
5.2	Compact web interface	22
5.3	Short response time	22
5.4	RESTful API	22
5.5	Import old data	23
6	Design	25
6.1	Application Structure	25
6.2	Caching	28
6.3	REST interface	30
6.4	Querying	32
7	Implementation	35
7.1	Haskell	35
7.2	Yesod Web Framework	41
8	Evaluation	51
8.1	Termination Competition 2014	51
8.2	Future considerations	52
8.3	StarExec	53
9	Summary	55
	Glossary	57
	References	61
	Eidesstattliche Versicherung	63

Abstract

In 2014 the annual Termination Competition was held for the eleventh time since 2004. It was organized by members of the Leipzig University of Applied Sciences (HTWK). This thesis gives an overview about the event, the *Termination Community* as well as an introduction to *Termination of Term Rewriting*. It also discusses StarExec, the technical infrastructure used by the competition. In general this is also an introduction to the *Star-Exec-Presenter*, a web-based application to present the results of the competition.

2014 wurde die jährliche Termination Competition zum elften Mal seit 2004 ausgetragen. Sie wurde durch Angehörige der Hochschule für Technik, Wirtschaft und Kultur Leipzig (HTWK) organisiert. Diese Masterarbeit gibt einen Überblick über das Ereignis, der *Terminations Community*, sowie eine Einführung in das Thema *Termination von Termersetzung*. Sie bespricht ebenso StarExec, die technische Infrastruktur, welche für den Wettbewerb genutzt wurde. Vor allem aber soll sie eine Einführung in *Star-Exec-Presenter* sein, das eine web-basierte Anwendung zur Präsentation der Resultate aus dem Wettbewerb ist.

Preface

With this thesis I hope to successfully graduate with the master's degree and therefore to end a journey I started over three years ago. My goal was to become a better programmer by deepening my theoretical knowledge and I am thankful for what I was able to achieve.

I decided to write this thesis with Prof. Dr. Waldmann as my advisor, because I wanted to have a challenge. I wanted to learn something new and this is exactly what I accomplished. I learned the programming language Haskell, which is not that easy if you're absolutely not experienced with functional programming. And I got in touch with server-side web development, which I always avoided due to my own lack of interest for PHP. I'm happy to have build new expertises and I will hopefully apply those in the future.

I want to thank my advisor Prof. Dr. Johannes Waldmann for giving me the chance to do this project but also for his always illuminating courses and his kindness. Without him I might have never programmed real Haskell code. I also want to thank Dr. René Thiemann from the University of Innsbruck, who is the second reviewer of this thesis besides Prof. Waldmann. Most notably, I want to thank my beautiful girlfriend Anne Wiegand for her love and support. I also want to thank my friends and colleagues, especially from PLUSPOL interactive, for the inspiring conversations regarding my thesis: Andreas Kießling, Ulrich Merkel, Thomas Stangneth, Christoph Schreiber and Martin Hentschel, just to name a few. Of course, I want to thank the CEOs of PLUSPOL Jörg Brücker, Stefan Dittmar and Thomas Lange for giving me the chance to make my graduation possible. Finally, I want to thank my sister Dr. Kerstin von der Krone for her guidance in the final state of this paper.

1

Introduction

This is a guide to the Star-Exec-Presenter, a web-based visualization and data-mining software developed in preparation of this thesis. The application is intended to be an efficient to use software for members of the Termination Community of the StarExec service, as well as other communities or users of StarExec. It features essential functionalities such as starting competitions, loading results and other data from StarExec or visualizing and filtering such data. Contributors of solvers should be able to track the progress of development of their respective solver by comparing its results over time.

The motivation of this thesis is to provide a tool for StarExec users for further research as well as some kind of standardization of hosting and running future Termination Competition and competitions of other StarExec communities. The goal for Star-Exec-Presenter is to be a tool, that is easy to use as well as easy to install and run. It is open source, so it can be forked und changed to better meet future needs.

Termination is a branch of theoretical computer science where programs such as algorithms are examined, whether they terminate, i.e. whether they complete. In particular, the termination of Term Rewriting System is of special interest for the Termination Community. A Term Rewriting System is a turing complete computational model such as Turing machines. In terms of StarExec, solvers are running over a bunch of problems returning an answer for each problem either with *YES*, *NO* or *MAYBE*.

The Termination Competition is an annual event of the termination community of StarExec, where solvers from different contributors compete in several categories. In 2014 the competition was organized by my adviser, Prof. Dr. Johannes Waldman, and me. Our goal was to do both, starting the competition run as well as to present it via an automatically updated web-interface.

The 2014 competition generated a total amount of 13 Gigabyte of data and used about 90 days of CPU time on all of the 193 Quad-Core CPUs of the StarExec cluster.¹ It took over eleven hours to complete all jobs.

Our additional goal was to be able to compare the results with those from previous competitions. So we added an option to import the data from the 2007 competition hosted by the Laboratory of Computer Science at the University of Paris-Sud and the following competitions hosted by the University of Innsbruck.

“*StarExec is a cross community logic solving service developed at the University of Iowa (...).*”² It is a technical infrastructure providing the service to run logic solvers on a powerful cluster of CPUs. It also provides an extensive web-based user interface to upload and run solvers and problems, referred to as benchmarks. A REST-based³ API⁴ is utilized by the Star-Exec-Presenter.

The Star-Exec-Presenter is the software developed as the concrete work for this thesis. It is a web application programmed in Haskell⁵ with the utilization of the Yesod Web Framework⁶.

In the following chapters I will go through the various topics of this thesis starting with the *Termination* where I will try to discuss *Termination of Term Rewriting*, the community itself and the *Termination Competition* as well as the *Termination Problem Data Base*, the compilation of problems which all participating tools have to solve. I will give an introduction to *StarExec* and how it works in the perspective of our tool, the Star-Exec-Presenter. I will list some *Use Cases* which we want to make possible. Afterwards these use cases are derived from by the *Requirements* which define what our tools has to be capable of. The *Design* and *Implementation* will then explain the internals of Star-Exec-Presenter by going into the details of its *Data Model* or the *Caching* mechanism as well as the actual code by giving an introduction to *Haskell* and the

¹see <http://lists.lri.fr/pipermail/termtools/2014-July/000979.html>

²see <https://www.starexec.org/starexec/public/about.jsp>

³Representational state transfer, a paradigm to implement a server-client communication

⁴Application programming interface, a set of components (methods, protocols, tools) for using a software

⁵a functional, non-strict, declarative programming language (<http://www.haskell.org/>)

⁶a REST-based web application framework (<http://www.yesodweb.com/>)

Yesod Web Framework. Finally, this thesis will be concluded by an *Evaluation*.

The complete source code of the Star-Exec-Presenter can be found on the CD appended to this thesis. It is also hosted on GitHub⁷. The latest version can be found in the `develop` branch.⁸

⁷see <https://github.com/stefanvonderkrone/star-exec-presenter>

⁸see <https://github.com/stefanvonderkrone/star-exec-presenter/tree/develop>

2

Termination

In this chapter I will discuss the terms of Termination and Termination of Term Rewriting. I also will describe the Termination Competition, as its 2014 incarnation was one of the goals towards the Star-Exec-Presenter was developed.

2.1 Termination of Term Rewriting

The motivation of this topic is to proof wether a term rewriting system (TRS) is terminating. Termination in terms of computer science means that a programm, e.g. a function, ends with a base condition. If this is the case then the program terminates. A base condition for example is the end of a list or counting down an integer and then reaching zero. A good example is the `replicate` function which takes an integer and an arbitrary argument which will be replicated into a list. The integer defines the amount of occurrences of that second argument in the resulting list. As this thesis also covers this programming language the following code listing is written in Haskell:

```

1 -- takes an Int and an a, returns a list of type a
2 replicate :: Int -> a -> [a]
3 -- the base condition, returns empty list if Int is zero
4 replicate 0 x = []
5 -- recursive call by decrementing n
6 replicate n x = x : replicate (n-1) x

```

Assuming, that we have a strict type system – which is the case with Haskell – and the first argument is a positive integer, then `replicate` always terminates. But, if the first argument is a negative integer, then this function will never ever terminate, because the base condition will never be fulfilled. So, this implementation of the `replicate` function is a good example of both, a terminating as well as a non-termination program. That’s why the actual implementation in Haskell’s base library is better.¹

A TRS itself is a set of rules each being a pair of terms $l \rightarrow r$ where the left term can be replaced by the right term, but not vice versa. A TRS is an alternative computational model to a Turing machine. A TRS terminates if there is no rule left to be applied. In the example function from above, there are two rules, i.e. two implementations of that function. If the first argument is zero, the call to this function will be replaced by an empty list. If this argument isn’t zero, the second rule will be applied. Then, the call to this function will be replaced by the second argument being appended to the result of another call to `replicate`, this time with `n` being decremented.

So, the question is, how can the termination of a TRS be proven? According to (Zantema 2000), there are *semantical*, *synthactical* and *transformational* methods for a proof.

“Proving termination by a semantical method means that a weight function has to be defined in such a way that by every reduction step the weight of a term strictly decreases. If the weight is a natural number, or more generally, a value in a set equipped with a well-founded order, then this cannot go on forever so termination is ensured.”

To give an example for a semantical method, a TRS with the following set of rules is given:

$$R = \begin{cases} f(f(x)) \rightarrow g(x) \\ g(g(x)) \rightarrow f(x) \end{cases}$$

¹see <http://hackage.haskell.org/package/base-4.7.0.2/docs/src/GHC-List.html#replicate>

Now, well-formed weight functions must be chosen or interpreted from the set of rules. So, we define these functions in terms of the non-empty set of A : $f_A(x) = g_A(x) = x + 1$ for all $x \in A$. These weight function can now be applied to the rules:

$$f_A(f_A(x)) = x + 2 > x + 1 = g_A(x) \text{ and } g_A(g_A(x)) = x + 2 > x + 1 = g_A(x)$$

The termination of this TRS has been proven, because the application of the weight functions leads to a decrementation of the weight of each term. This example is taken from (Zantema 2000) and it is just one of a terminating TRS proven by sementical methods. The other two methods are more complex and explained as follows:

“Syntactical methods are based upon orders on terms that are defined by induction on the structure of the terms. Given such an order, if it can be proved that it is well-founded, and if every reduction step causes a decrease with respect to the order, then termination has been proved. (...) Transformational methods provide non-termination preserving transformations between rewrite systems: a transformation Φ falls in this category if termination of a TRS (Σ, R) follows from termination of $\Phi(\Sigma, R)$. If such a Φ can be found in such a way that termination of $\Phi(\Sigma, R)$ can be proved by any other method, then termination of (Σ, R) has been proven.” (Zantema 2000)

2.2 Termination Community

The Termination Community is an international group of developers and researchers working in the field of Termination of Term Rewriting. Their main hub of information is the website <http://www.termination-portal.org/>, a wiki containing all sorts of information about this community. They organize workshops where they try to meet up with other researchers interested in all facets of termination.² At these workshops they discuss new scientific achievements in termination. The first meeting was in 1993 while the latest one was held during the Vienna Summer of Logic where also the latest Termination Competition was officially presented. The Termination Community also has a mailing list³ where they discuss all sorts of topics regarding termination. Research groups on termination exist at the RWTH Aachen University, the Eindhoven University of Technology, the Institute of Computer Science at the University of Innsbruck, the Leipzig University of Applied Sciences and the Computer Research Laboratory at the University of Paris-Sud.

²see <http://www.termination-portal.org/wiki/WST>

³see <http://lists.lri.fr/mailman/listinfo/termtools>

2.3 Termination Competition

The competition is an event where several programs compete against each other solving problems regarding the Termination of Term Rewriting. There are many different categories with varying requirements.⁴ But only those categories represented by a significant number of tools take part at the official competition run. There are two types of scoring by which each tool is evaluated. The first one considers the number of results being *YES*, the second is the sum of an additional complexity degree appended to *YES*. So, only if a tool outputs the result *YES* regarding a certain problem it is rewarded with a score. Another score relates to meta-categories, groups of categories, where solvers are classified in an overall ranking. This score is calculated by the position in each category. The resulting score for each tool is the sum of points rewarded in each category. These points are the result of the number of all participants subtracted by the rank of the tool in the particular category.

The Termination Competition is an event held every year since 2005. The idea of comparing tools for proving term rewriting was brought up by Albert Rubio in 2003 (Marché and Zantema 2007). Until 2007 the competition was organized by Claude Marché of the Computer Research Laboratory at the University of Paris-Sud. From 2008 until 2013 the organization was in control of René Thiemann of the Institute of Computer Science at the University of Innsbruck. In 2014 Prof. Dr. Johannes Waldmann was in charge of the Termination Competition.

The Termination Competition is controlled by the steering committee, a group of researchers representing all research groups of the Termination Community. This committee influences the design and the process of each competition and it is responsible for the categories, the scoring as well as the problem database. The Termination Competition Steering Committee Bylaws⁵ define all rules by which the committee operates. By the time of writing of this thesis the committee is staffed by Jürgen Giesl, Frederic Mesnard, Albert Rubio, René Thiemann and Johannes Waldmann. Besides the steering committee there is also an organization committee which is responsible for the staging of a competition. The Termination Competition 2014 was organized mainly by my advisor Prof. Dr. Waldmann with my assistance in developing the Star-Exec-Presenter.

⁴see <http://www.termination-portal.org/wiki/Category:Categories>

⁵see http://www.termination-portal.org/wiki/Termination_Competition_Steering_Committee_Bylaws

2.4 TPDB - Termination Problem Data Base

The Termination Problem Data Base is a large collection of termination problems and the crucial part of each competition besides the participating tools. They contain challenging examples by various authors in several categories. Since 2009 all problems are described in XML format, until 2008 the problems were described in a special text based format. The version used for the 2014 Termination Competition was 9.0.⁶

⁶see <http://www.termination-portal.org/wiki/TPDB>

3

StarExec

StarExec is the service that is utilized by the Termination Community and other communities to run their research as well as the latest Termination Competition. This chapter will examine the interfaces from which the relevant data is fetched as well as their special characteristics. I will go into the details of how StarExec works and how to get the data.

3.1 History

StarExec is *“a public web-based service built to facilitate the experimental evolution of logic solvers, broadly understood as automated tools based on formal reasoning. (...) The service, (...), is designed to provide a single piece of storage and computing infrastructure to logic solving communities and their members. It aims at reducing duplication of effort and resources as well as enabling individual researchers or groups with no access to comparable infrastructure.”* (Stump, Sutcliffe, and Tinelli 2014) The work on StarExec started in 2012 as a joint development of the University of Iowa and the University of Miami. It is a cluster of 193 quad-core CPUs and 128 to 256 Gigabyte main memory each. The whole system is managed by a Red Hat Enterprise Linux.

The first public event was held in summer 2013. One year later, the Termination Competition 2014 was realized with StarExec.

3.2 Data Model

The data model has three major entities: the jobs, the solvers and the benchmarks. A benchmark is a formular¹ regarding a formal problem which will be processed by a solver. A solver is a programm that takes a benchmark, analyses it and gives an output regarding the given task. For example, in the case of the Termination Competition, this output should mainly consist of one of the following results: *YES*, *NO*, *MAYBE*. Whereby the answer *YES* can also contain additional information for the complexity of the problem. A post-processor is responsible for filtering this info out of the solvers log.

A job is a complex task with a set of solver and benchmarks, where each solver processes each benchmark. Each unique relation between a certain solver and benchmark is called *job-pair*. Each job-pair consists of the solver's log (stdout and stderr). The job-pair is directly connected to the result info, which contains the solver's result and additional info such as cpu- and wallclock-time, the configuration and identifying information of the certain solver and benchmark.

So, when a job is started, each solver will be started as well with a benchmark as input. The StarExec infrastructure itself takes care that all available CPUs are busy. Of course, not all job-pairs are started at the same time. If a solver finished its task a post-processor is working on the solver's output to extract the actual result (*YES*, *NO*, *MAYBE*). All information will then be stored in the MySQL database. The following figure illustrates this process:

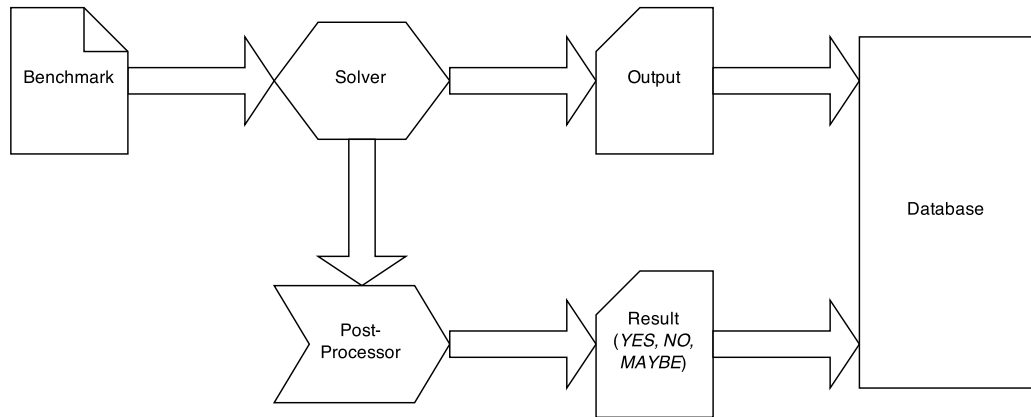


Figure 3.1: the process of a benchmark

Hierarchically, StarExec is structured into communities which use spaces for further subdivision. Currently, there are several communities like *Termination* or *SMT*. Each

¹see <https://wiki.uiowa.edu/display/stardev/Project+Overview>

community consists of several spaces, some being owned by a user, some being public. Each space can also have subspaces. A space consists of a bunch of jobs, solvers and benchmarks. The job-results and job-pairs can be accessed through the particular job. Solvers and benchmarks can be used in other spaces by copying or linking them.

3.3 Interfaces

StarExec has an extensive interface to communicate with. It provides a complex user interface via the browser as well as a RESTlike.

3.3.1 GUI

The browser based web interface of StarExec is characterized by the hierarchical nature of the spaces. A user has to have access to a space to start a job as well as upload a solver or benchmark. Each space has an overview of its job, solvers, benchmarks, its subspaces and even the enlisted users. All entities listed in this overview are owned by the particular space, so the entities of a subspace can only be seen by navigating to the certain space.

As addressed before solvers and benchmarks can be copied or linked from other spaces but also be uploaded into the space. The available actions may alter according to how much access a user may have regarding a certain space. Some space may even be hidden from the user. Aside from starting jobs and other actions mentioned before, it is also possible to download certain information such as all jobs and benchmarks or a XML based representation of the space itself.

The web interface relies heavily on AJAX-request² to prevent each page from being reloaded on each action. So, when the user is navigating the tree of spaces, there is no actual reload of the whole page but an asynchronous fetching and updating of the page's content. This behavior is also applied to the handling of a job's progress. So, until a job is finished the web page automatically keeps track of the progress by regularly sending AJAX-requests to StarExec.

²Asynchronous JavaScript And XML, a technique to dynamically load data remotely and update the content

3.3.2 API

Another part of the interface to StarExec is also a part of the GUI but only within download-links or forms: a RESTlike API. This interface is heavily used by the java tool StarExecCommand³ which is a commandline tool for the communication with StarExec. It provides a shell-like prompt that can be used to do certain action on StarExec. StarExecCommand was also some kind of inspiration for the design of Star-Exec-Presenter when it comes to the interface to StarExec.

In the following I want to list the important URLs that are also used by the Star-Exec-Presenter:

https://www.starexec.org/starexec/secure/j_security_check

Of course, without this URL StarExec is actually of no use, because a user has to be logged in. With the correct credentials, a POST request has to be send to StarExec. If StarExec returns a valid session ID the user is successfully logged in. Star-Exec-Presenter itself doesn't require the client to provide his own StarExec credentials, as the application uses its own. So, everyone who starts our web application has to provide valid login credentials.

<https://www.starexec.org/starexec/secure/download>

Another important URL is the one to download certain data. This URL is used consistently throughout all download links. It requires GET parameters attached to it, the most important being the `id` and the `type`. The parameter `type` defines the type of entity such as `job` or `spaceXML`. Whereas the parameter `id` has to be a valid identifier of the entity. If `type` is set to `job` the job-results will be downloaded in the form of a zipped CSV. Alternatively, if `type` is set to `spaceXML` a zipped XML containing the space's hierarchy will be downloaded.

<https://www.starexec.org/starexec/services/jobs/pairs/{pairId}/stdout>

<https://www.starexec.org/starexec/services/jobs/pairs/{pairId}/log>

Both listed URLs are related to the job-results. As the job-results consist only of some meta-data as well as the final result, the actual output has to be downloaded via these two URLs. Conveniently, the response is plain text. The placeholder `{pairId}` refers to the identifier of a job-result.

³see <https://www.starexec.org/starexec/public/starexeccommand.jsp>

<https://www.starexec.org/starexec/secure/details/job.jsp>
<https://www.starexec.org/starexec/secure/details/solver.jsp>
<https://www.starexec.org/starexec/secure/details/benchmark.jsp>
<https://www.starexec.org/starexec/secure/edit/processor.jsp>

To receive additional infos of the certain entities with which Star-Exec-Presenter works, there are four additional URLs to observe. As StarExec has no way to get these information other than via web pages, Star-Exec-Presenter parses the required data from HTML. Each URL requires a GET parameter id set to the particular entity's identifier attached to the end.

4

Use Cases

In this chapter I will list the use cases which the Star-Exec-Presenter is mapped onto. These use cases range from simple tasks, for instance displaying solver-results from StarExec, up to more complex ones, like starting and monitoring a competition.

4.1 Start a predefined Competition

One of the core use cases of the Star-Exec-Presenter is the ability to start a competition run on StarExec. As StarExec actually doesn't know what a competition in our terms is, Star-Exec-Presenter should know it. So all meta-information of a competition should be managed by the Star-Exec-Presenter. These meta information include hierarchical information of the competition's organisation as well as the details which we'll send to or get from StarExec.

Hierarchical information is meant as how each solver and benchmark is organised within the competition. Each competition consists of several meta categories which cover a general topic in the competition. The 2014 Termination Competition has the following categories: *Termination of Term Rewriting (and Transition Systems)*, *Complexity Analysis of Term Rewriting* and *Termination of Programming Languages*.

Each meta category is divided into smaller, more specific categories, each being a representation of an actual job on StarExec. These categories may be best described

by the meta categories *Termination of Programming Languages*, as there is a category about the programming language C. Other categories could bring up Java, Logic Programming or Functional Programming.

4.2 Displaying Results and Outputs

The core function of StarExec is not only to run the solver on specific benchmarks but also to host the results and outputs of each run. So a major task for the StarExec-Presenter is to display that data in an appropriate way. The explicit data is the result of a solver with a specific benchmark and its CPU- as well as Wallclock-Time. Additionally the solver's output is important which is accessible on StarExec with the job-pair, a unique combination of solver and benchmark within the specific job. The output can contain plain text or even XML or HTML. Figure 4.1 is a screenshot of the results of a job.

Overview of job-results

[flexible query \(experimental\)](#) | view original jobs on star-exec:[5374](#),

Statistics

5260 pairs, 1078452.9 / 683824.1 s

Results

[Show Legend](#)

Jobs	SRS Stand 37909			
Solver	NaTT	TTT2	mu-term 5.13	AProVE_JRE2
Scores	204	607	139	839
SRS_Standard/Bouchare_06/01.xml	0.0 / 0.2 s	36.7 / 9.6 s	299.9 / 300.0 s	6.2 / 1.9 s
SRS_Standard/Bouchare_06/02.xml	0.0 / 0.1 s	3.1 / 1.0 s	299.2 / 300.0 s	6.8 / 2.1 s
SRS_Standard/Bouchare_06/03.xml	0.0 / 0.2 s	27.6 / 7.2 s	299.9 / 300.0 s	38.0 / 10.1 s
SRS_Standard/Bouchare_06/04.xml	0.0 / 0.2 s	38.8 / 9.8 s	299.9 / 300.0 s	5.3 / 1.7 s
SRS_Standard/Bouchare_06/05.xml	0.0 / 0.1 s	36.1 / 9.4 s	300.0 / 300.0 s	48.1 / 12.8 s
SRS_Standard/Bouchare_06/06.xml	0.0 / 0.5 s	32.3 / 8.4 s	299.9 / 300.0 s	32.3 / 8.9 s
SRS_Standard/Bouchare_06/07.xml	0.0 / 0.1 s	31.0 / 8.1 s	2.0 / 2.1 s	52.4 / 14.0 s
SRS_Standard/Bouchare_06/08.xml	0.0 / 0.2 s	0.7 / 0.4 s	298.5 / 300.1 s	1.5 / 0.7 s

Figure 4.1: example of an overview of job-results

4.2.1 Displaying a Competition and its Results

The Star-Exec-Presenter is a tool that ran during the 2014 Termination Competition. It should display the current status of the competition run with a compact web interface that automatically updates itself. Figure 4.2 shows a screenshot of the implemented competition overview.

Termination Competition 2014

General Information wc = 300 a = 100 b = 1000 c = 1.0 (2014-07-20 02:25:05.328531 UTC) 37880 pairs, 7806889.5 / 3143466.2 s finished in 11h 41m 27s

Termination of Term Rewriting (and Transition Systems) finished in 11h 41m 27s, 27599 pairs, 4280960.8 / 2004985.2 s

Combined Ranking (Rules): 1. AProVE_JRE2 (19) 2. TTT2 (5) 3. NaTT (4) 4. CppInV (3) mu-term 5.13 (3) 6. T2 - 2014-07-06v1 (2)
7. matchbox2014.07.08 (1) Wanda (1) 9. THOR (0) Ctrl (0)

category	post-proc	rankings	statistics
TRS Standard	plain.2	AProVE_JRE2 (1312), NaTT (1024), TTT2 (997), mu-term 5.13 (856), Wanda (636),	7405 pairs, 416831.1 / 247239.1 s
SRS Standard	plain.2	AProVE_JRE2 (839), TTT2 (607), NaTT (204), mu-term 5.13 (139),	5260 pairs, 1078452.9 / 683824.1 s
TRS Relative	plain.2	AProVE_JRE2 (35), TTT2 (24),	88 pairs, 13062.8 / 4027.5 s
SRS Relative	plain.2	AProVE_JRE2 (89), TTT2 (25),	410 pairs, 106678.3 / 45449.7 s
TRS Standard certified	ceta.9	AProVE_JRE2 (1222), TTT2 (966), matchbox2014.07.08 (714),	4443 pairs, 774987.2 / 331790.1 s
SRS Standard certified	ceta.9	AProVE_JRE2 (819), matchbox2014.07.08 (584), TTT2 (572),	3945 pairs, 1465907.9 / 437506.9 s
TRS Relative certified	ceta.9	AProVE_JRE2 (31), TTT2 (24),	88 pairs, 13971.7 / 4898.0 s
SRS Relative certified	ceta.9	AProVE_JRE2 (89), TTT2 (23),	410 pairs, 96316.9 / 45322.1 s
TRS Equational	plain.2	AProVE_JRE2 (62), mu-term 5.13 (59),	142 pairs, 2898.8 / 2661.4 s
TRS Conditional	plain.2	mu-term 5.13 (25), AProVE_JRE2 (18),	64 pairs, 4401.4 / 2887.3 s
TRS Context Sensitive	plain.2	mu-term 5.13 (101), AProVE_JRE2 (98),	216 pairs, 8366.2 / 5349.0 s
TRS Innermost	plain.2	AProVE_JRE2 (273), mu-term 5.13 (211),	732 pairs, 124684.1 / 75201.2 s
Higher-Order rewriting (union beta)	plain.2	Wanda (152), THOR (121),	350 pairs, 3875.1 / 1409.1 s
Integer Transition Systems	plain.2	CppInV (747), T2 - 2014-07-06v1 (562), AProVE_JRE2 (217), Ctrl (184),	3812 pairs, 156180.8 / 112903.0 s
Integer TRS	plain.2	AProVE_JRE2 (104), Ctrl (80),	234 pairs, 14345.7 / 4516.8 s

Complexity Analysis of Term Rewriting finished in 10h 38m 43s, 8988 pairs, 3453948.0 / 1095372.9 s

Combined Ranking (Rules): 1. TCT (2) 2. AProVE_JRE2 (1) 3. CaT (0)

category	post-proc	rankings	statistics
Derivational Complexity - Full Rewriting	plain.2	TCT (1575), CaT (1108),	3618 pairs, 1611245.4 / 474675.2 s
Runtime Complexity - Full Rewriting	plain.2	TCT (1505), CaT (889),	2578 pairs, 722622.5 / 236598.8 s
Runtime Complexity - Innermost Rewriting	plain.2	AProVE_JRE2 (2179), TCT (2082),	2792 pairs, 1120080.1 / 384098.9 s

Termination of Programming Languages finished in 5h 26m 50s, 1293 pairs, 71980.7 / 43108.1 s

Combined Ranking (Rules): 1. AProVE_JRE2 (2) 2. UltimateBuchiAutomizer (1) 3. T2 - 2014-07-06v1 (0)

category	post-proc	rankings	statistics
C	plain.2	AProVE_JRE2 (296), UltimateBuchiAutomizer (271), T2 - 2014-07-06v1 (183),	1293 pairs, 71980.7 / 43108.1 s

General Information

copyright: Stefan von der Krone, Johannes Waldmann | F-IMN | HTWK Leipzig

Figure 4.2: overview of the Termination Competition 2014

4.2.2 Interpret Outputs as HTML or XML

Some solver explicitly log an HTML- or XML-proof with each run in its output. This proof should be detected as well as extracted from the output. Also the respective XSL-File for the XML output should be used to display the proof in the browser.

4.3 Compare Results, Querying

Simply displaying the results and outputs of each solver-run is not enough. It is more important to compare the results between several solver. Are there differences between their results on specific benchmarks? Also, a list of jobs should be displayable at once. A historical comparison is the most important aspect of this use case, as the developers of a solver would want to track the evolution of their tools.

As important comparing the result is as overwhelming is their amount, so a beneficial use case is querying the results according to specific filters. Such filters can eliminate unimportant results from the representation. For instance a benchmark is uninteresting if all solvers have the same result, so all those benchmarks can be filtered out.

The user should be able to concentrate on data which meets his or her interests. This data could tell which benchmark causes incorrect results with the own solver. Also it can illustrate changes in the solvers results in comparison to past runs, fo instance a once solved benchmark could now be unsolvable. This is the data we want to query in terms of data mining.

5

Requirements

This chapter discusses the requirements that were agreed to. Which data is interesting and how is it handled? How should the application be developed, which kind of data should be accessible through the Star-Exec-Presenter? And how should the application incorporate with StarExec? All requirements are inferred from the previous chapter.

5.1 Solver and Post processor output

To display the results of a job on StarExec we must know which data actually is important. So, basically StarExec knows two different kinds of resulting data. The first is a table which contains the resulting data of each solver-benchmark combination (the job-pair) of the job. It tells us the final result of a job-pair as well as its CPU- or Wallclock-Time and its current status. Each job-pair has a unique identifier which leads us to the second kind of resulting data: the output of a job-pair which is accessible separately. It contains the standard-out of the solver and a log generated by StarExec.

The aforesaid table should be displayed wellarranged with the Star-Exec-Presenter. Important data kinds are of course the result of the job-pair as well as the CPU- and Wallclock-Time. The presentation should be easy to read and its information should

be gathered quickly.

The output of the job-pair should be scanned for XML to determine whether it contains a proof. The corresponding presentation should give the chance to display that proof according to its required XSL-Stylesheet.

All data should be persistently stored in a Postgre-SQL database.

5.2 Compact web interface

The web interface of StarExec is quite verbose and not specifically matching the needs of members of the termination community. So, the Star-Exec-Presenter should feature a much more compact web interface which clearly guides the user to the destined information. Readability is important.

5.3 Short response time

As StarExec is located in Iowa and some requests require expensive fetching from its database some responses may need up to seconds to be fully available. This is not desired as some users may consider it as an error of the Star-Exec-Presenter. So our tool should cache all data that it loads from StarExec. An additional information should be shown when such data is about to be fetched. The desired effect is that each request is responded as early as possible. Also every additional calculation which is repeatedly needed should be cached, too.

5.4 RESTful API

The Termination Community is filled with members living all over the world and they may want to be able to send interesting data by e-mail. And as it is not necessary to pack all data into one e-mail, why not simply send a URL? So, another requirement should be to have distinctive URLs which contain all informationen required to display the desired data. Such a URL can point to the results of a job on StarExec and only requires the identifier of that particular job.

5.5 Import old data

There is one use case which defines that a user should be able to compare certain job results. These results may not only originate from a job on StarExec but also from previous competitions whose data may not be available on StarExec. So, Star-Exec-Presenter should be able to import as well as to handle such data.

6

Design

The Star-Exec-Presenter is actually the concrete work of this thesis. As a REST-based web application it is an interface to the StarExec service and it provides a caching mechanism. This chapter will give detailed insights in the architectural approach to build the Star-Exec-Presenter according to the requirements.

6.1 Application Structure

Star-Exec-Presenter is written entirely in Haskell by utilizing its module system and the Yesod Web Framework. Detailed information are listed in chapter 7. *Implementation*. In this section I will discuss the particular pieces that form Star-Exec-Presenter into a full server-side web application. To better understand the application structure, it could easily be interpreted as a Model-View-Controller (MVC) pattern where the Model stands for the data model, the View is represented by the templates which form the HTML presentation, and the Controller in the form of the several request handlers.

6.1.1 Data Model

The data model can be divided into two groups, the first containing the data types used for the REST-API which I will discuss later, and the second being the persistent data types of the actual data like the job results.

The first group of data types is used as identifiers for the second group, so they are directly related to each other. For instance there is a `JobID` type that relates to `Job` type that contains the infos about a job as well as to a list of `JobResults` containing the actual results of the particular job. All these types mostly are unified types for the different entities of `StarExec` as well as imported entities.

To complete the example from above the type `JobID` has three different appearances, first being `StarExecJobID` representing a job on `StarExec`, the second being `LriJobID` representing an imported job from the 2007 Termination Competition and the third being `UibkJobID` for the imported jobs from Innsbruck. Additionally, a `Job` also has three different appearances (`StarExecJob`, `LriJob`, `UibkJob`) as well as `JobResult` (`StarExecResult`, `LriResult`, `UibkResult`). The Figure 6.1 illustrates these dependencies.

The reason to have three different kinds of appearances for each type is because of having the data from `StarExec` in combination with the imported data from previous competitions. This old data differs in its form from that of the 2014 Termination Competition, so to consider these differences in a safe way I decided to isolate them persistent-wise. So, for each kind of data there are three database tables used to store that data.

Code-wise, the data model can be found within the `Presenter.Model` module, which combines all the groups of data types mentioned above as well as the persistent data types, which are defined in the `config/models` file.

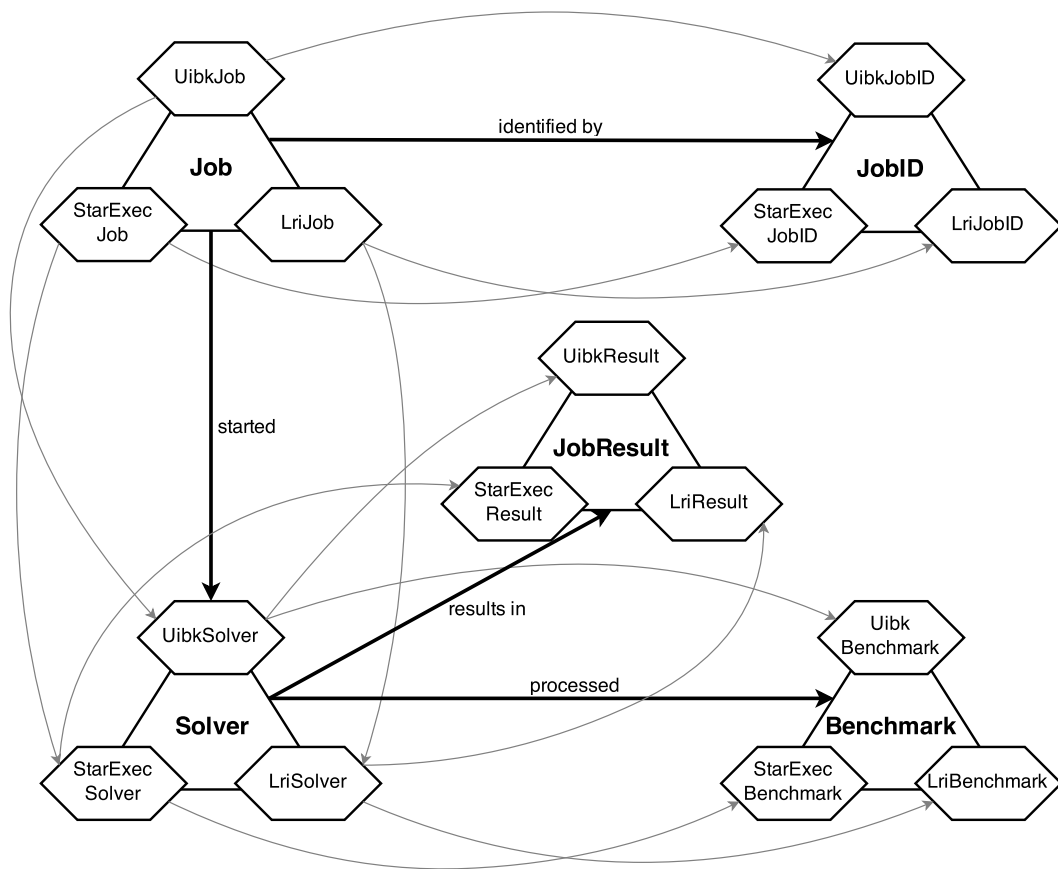


Figure 6.1: dependencies of essential entities

6.1.2 Templates

The templates are essentially the view of Star-Exec-Presenter, they are filled with the requested information in a proper way to be displayed via the web browser. HTML¹ is very well suited for this purpose. Other types of data formats like JSON², CSV³ or XML⁴ can be considered as well but aren't a subject of matter, although they are used in the communication with the StarExec service.

All templates can be found in the `templates` folder, some are located within the `Presenter.Utils` module.

6.1.3 Request Handlers

The controller part of Star-Exec-Presenter is represented by the several request handlers. Each handler responds to a certain URL⁵ requested by a client application. For instance there are handlers that only return the information of a particular entity like a solver or a job, there are also handlers that initiate specific requests to StarExec or display the results of a competition. The request handlers are generally managing the logic behind Star-Exec-Presenter.

All handlers are implemented within the `Handler` module.

6.2 Caching

Fast responses, especially for handlers that can return a huge bunch of data, is a main requirement for Star-Exec-Presenter, so there has to be an effective solution. We came up with the simple idea of caching all data in the database. For instance, when the user requests the results of one or more certain jobs, the application first looks in the database for the respective data and returns it. If there isn't found any related data it will be fetched from StarExec. But this fetching can affect a fast response, so it will be separated from the actual response by running in the background.

This is achieved by starting a new thread within the Star-Exec-Presenter application. The following figure shows the process from a request to a response:

¹Hypertext Markup Language, a markup language designed to be interpreted by web browsers

²Javascript Object Notation, a compact subset of javascript designed for data exchange

³Comma-Separated Values, a data format to represent tables or list

⁴Extensible Markup Language, a data format for hierarchically structured data like trees

⁵Uniform Resource Locator

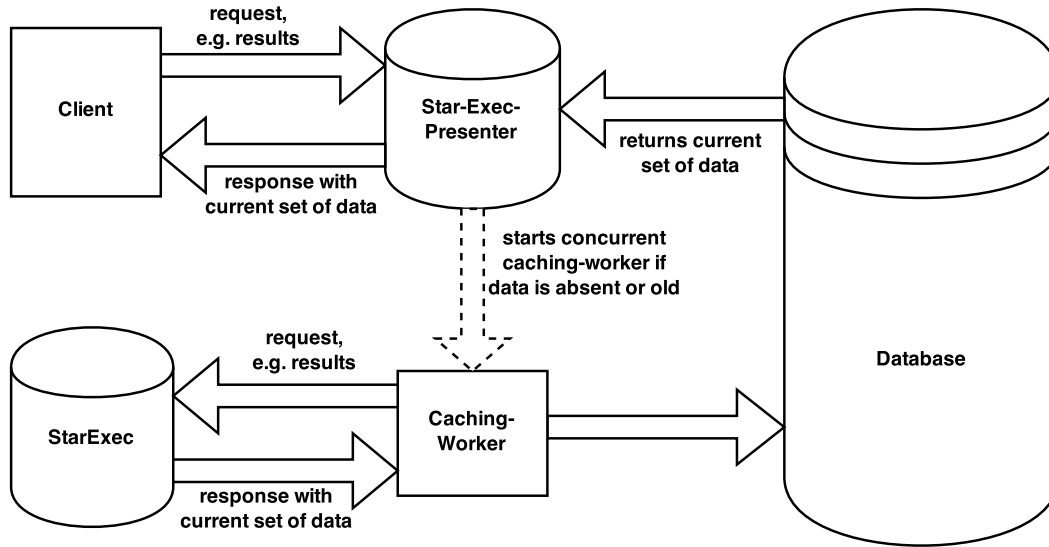


Figure 6.2: the caching mechanism

This database based caching mechanism has its downside when it comes the the competition results, as they are calculated based on the results of the related jobs. And as the data of these jobs is stored in the database, there is no need for the results to be stored either. So the results have to be cached in another way which is within the application's memory. For this purpose Star-Exec-Presenter uses STM⁶.

The reason to cache the competition results is that it takes too long to fetch the related job results from the database and do the calculations, as a usual competition has a huge amount of data to process. For instance the 2014 Termination Competition has three meta categories with a total of 19 jobs that produced an overall amount of 37.880 job-pairs. That doesn't seem to be that much but one requirement was to have short response times. So we decided to offload the computation to a separate worker thread. The results are saved within the application's memory and are accessed via STM. The calculating worker thread is the only instance which writes in this cache, all request handlers can only read from it.

STM helps in this case because it can prevent deadlocks from multiple simultaneous accesses to the shared memory making the whole application more reliable. To achieve this goal every access to the shared memory (our cache of competition results) is being managed by transactions, much like transactions of a database. If two or more transactions try to access a resource simultaneously, all but one will be

⁶Software Transactional Memory, a mechanism for concurrent programming with shared resources

aborted. A mechanism to retry a rolled back transaction ensures that they will be finished anyhow. To achieve this behavior each transaction must be very small. The actual terminology is *atomic*. It is preferential to have very small transactions because STM is not *fair* as it favors small and fast transactions over large and slow ones. STM doesn't work in FIFO⁷ order. (Marlow 2013)

Star-Exec-Presenter itself only does the reading and writing process within STM. The calculation of the competition results is done outside of STM. Only in this way we can ensure that the transactions are very small – or *atomic*.

6.3 REST interface

Star-Exec-Presenter is a server-side web application which receives requests and responses with an appropriate HTML presentation. These requests have to point to a valid resource, a URL. A list of important resource-URLs of the application including their accepted HTTP methods follows:

1	/jobs/#JobID	GET
2	/solvers/#SolverID	GET
3	/benchmarks/#BenchmarkID	GET
4	/pairs/#JobPairID	GET
5	/post_procs/#PostProcID	GET
6	/results/#Query/*JobIds	GET
7	/proofs/#Text	GET
8	/registered	GET
9	/competitions/#CompetitionInfoId	GET
10	/control	GET POST
11	/import	GET POST

Of course, an explanation of these routes, as they are called, is helpful. The first five routes are self-explaining. They take a specific identifier and return the infos of the related entity, whether it's a job, solver, benchmark, job-pair or post-processor. All these routes have a counterpart which lists all related entities. These routes are accessed without any parameter at the end.

The Results-Route takes a list of Job-Identifiers and returns all of their results. A special part of this route is its first dynamic path piece, the query parameter which

⁷First In First Out

defines a special case for this route. If there is given a specific query then the job results will be filtered according to it. More on this topic can be read in the section “Querying”.

The Proof-Route is directly related to the Job-Pair-Route. If a job-pair contains a proof, this route is meant to display it. The Registered-Route lists all of the participants of the 2014 Termination Competition including links to their configuration on StarExec. The Competition-Route displays the results for a specific competition. This route has a counterpart, too, which lists all (publicly available) competitions.

The Control-Route is an interface to start a competition whether it's a full or a small one. Requested with the default HTTP method GET this route only displays its interface form. If this form is submitted the POST method is requested and a competition will be started. The Import-Route is used to import old data from previous competitions. It has a small form with a select field and a file input to upload a file. It takes a zip file with the content which has to be defined by the select field.

There are other routes that give further information as well as some routes that are older versions of the aforementioned ones. They are still accessible for legacy purposes and redirect to the new ones.

As Star-Exec-Presenter is meant to be a simple interface for StarExec and the Termination Competition in the first place, these routes are quite simple as well. They only accept GET or POST requests and try to work according to principles of REST. REST is an acronym for Representational State Transfer and was first mentioned by Roy Fielding in his dissertation “Architectural Styles and the Design of Network-based Software Architectures”. REST respects the stateless nature of HTTP and defines that every resource has its unique uniform resource identifier (URI). (Fielding 2000)

HTTP methods like GET, POST, PUT or DELETE are used to get, create, change or remove the specific resource. The GET method is nullipotent, which means that a GET request never has any side effects and never changes anything respectively. POST, PUT and DELETE requests are idempotent, resulting in the same outcome no matter how much they are called. The URI, e.g. `/clear_the_database` itself never causes a change to the data unless it is requested with one of the mentioned HTTP methods. Apart from that, this is just an example which isn't recommendable for a real world REST API.

REST utilizes the available web technology for this purpose and isn't based on any standard like SOAP⁸. In contrast to SOAP REST is more like an architectural strategy or design pattern to develop a web application, whereas SOAP is a protocol. REST relies on the hypertext transfer protocol (HTTP) to mediate the request and simplifies the development process with it. To explain this principle a little further I want to emphasize the aforementioned MVC design pattern. In this context a REST interface can be considered the controller of the application which responds to the request with a suitable presentation.

A suitable presentation is HTML as well as XML or JSON. Star-Exec-Presenter currently only sends HTML. But this behavior could be altered by using a special header in the HTTP request. So, to get the data as XML or JSON, the Accept-Header could be set to the values `Accept: text/xml` or `Accept: text/json`. The web application then can evaluate this header to respond either with an XML or with a JSON presentation. (Richardson 2007)

It is widely accepted not to use verbs as parts of the URL rather than plural nouns. For instance the route `/get_job_results/#JobID` could be better written as `/job_results/#JobID` or simply `/results/#JobID`. Also, it is recommended to have a generalized route, e.g. `/jobs`, which lists all resources, whereas a parameter added to the URL results in returning only the specified resource. In the case of Star-Exec-Presenter especially the Results-Route is a special case that takes an arbitrary amount of JobIDs.⁹

6.4 Querying

Just presenting the results in a table of up to over 5000 rows is not very effective and can be a bit tedious, if the goal is to collect specific information. These kind of information can be in general the differences between each attending solver. The benchmarks that are completely – i.e. by all solvers – resolved to *YES* aren't interesting but the remaining ones. So, the most interesting benchmarks can be the ones, where one solver outputs *YES* while another one returns *NO*. For the developers this can indicate a bug or a lack of features.

Another interesting case of application is the historical comparison. So, if previous runs of one solver with the same set of benchmarks are available, then differing

⁸Simple Object Access Protocol, a web protocol for structured information

⁹e.g., see <http://www.sitepoint.com/best-practices-rest-api-scratch-introduction/>

results imply some kind of evolution to the solver. Of course, that evolution can be positive and negative as well.

The details of this querying are defined by the module `Presenter.Model.Query`. A query is defined by a list of transformations. A transformation can affect certain columns and rows. Columns are selected by their position within the table. Rows aren't directly selected but filtered by certain conditions. These conditions can include but also exclude specific rows. Each filter contains a list related to the list of solvers. For example, if the listing of results contains four solvers, then the filter contains four items. Currently, these items only enable a textbased comparison but no disjunction. Figure 6.3 shows a screenshot of the current implementation of the query interface.

By the time of writing of this thesis this definition is not finished and can only be considered as a draft. Another proposal is defined in the aforementioned module, too. This definition tries to be more generic in respect to the `SolverResult`. It also includes disjunction to the row-filter.

```
Query [Filter_Rows (And [Equals "nothing",Equals "solver=yes",Equals "solver=maybe",Equals "solver=yes",Equals "solver=yes"]
)]
```

summary

total number of rows: 1315

columns, by tags

Benchmark	NaTT/default (5374)	TTT2/ttt2 (5374)	mu-term 5.13/default (5374)	AProVE_JRE2/trs (5374)
nothing 1315 these others				
	solver-maybe 1111 these others	solver-maybe 708 these others	solver-maybe 1176 these others	solver-maybe 476 these others
		solver-no 37 these others	solver-no 1 these others	solver-no 110 these others
	solver=yes 204 these others	solver=yes 570 these others	solver=yes 138 these others	solver=yes 729 these others

row types

Benchmark	NaTT/default (5374)	TTT2/ttt2 (5374)	mu-term 5.13/default (5374)	AProVE_JRE2/trs (5374)			
nothing	solver-maybe	solver-maybe	solver=yes	solver=yes	1	these	others
nothing	solver-maybe	solver-no	solver-maybe	solver-maybe	1	these	others
nothing	solver-maybe	solver-no	solver-no	solver-no	1	these	others
nothing	solver=yes	solver-maybe	solver-maybe	solver-maybe	1	these	others
nothing	solver=yes	solver=yes	solver-maybe	solver-maybe	2	these	others
nothing	solver=yes	solver-maybe	solver=yes	solver=yes	8	these	others
nothing	solver-maybe	solver=yes	solver-maybe	solver-maybe	22	these	others
nothing	solver=yes	solver-maybe	solver-maybe	solver=yes	28	these	others
nothing	solver-maybe	solver=yes	solver=yes	solver=yes	30	these	others
nothing	solver-maybe	solver-no	solver-maybe	solver-no	35	these	others
nothing	solver=yes	solver=yes	solver-maybe	solver=yes	66	these	others
nothing	solver-maybe	solver-maybe	solver-maybe	solver-no	74	these	others
nothing	solver=yes	solver=yes	solver=yes	solver=yes	99	these	others
nothing	solver-maybe	solver-maybe	solver-maybe	solver=yes	146	these	others
nothing	solver-maybe	solver=yes	solver-maybe	solver=yes	351	these	others
nothing	solver-maybe	solver-maybe	solver-maybe	solver-maybe	450	these	others

[remove following 1 transformations](#)

apply transformation

```
Filter_Rows (And [Equals "nothing",Equals "solver=yes",Equals "solver=maybe",Equals "solver=yes",Equals "solver=yes"])
```

summary

total number of rows: 8

columns, by tags

Benchmark	NaTT/default (5374)	TTT2/ttt2 (5374)	mu-term 5.13/default (5374)	AProVE_JRE2/trs (5374)
nothing 8 these others				
		solver-maybe 8 these others		
	solver=yes 8 these others		solver=yes 8 these others	solver=yes 8 these others

row types

Benchmark	NaTT/default (5374)	TTT2/ttt2 (5374)	mu-term 5.13/default (5374)	AProVE_JRE2/trs (5374)			
nothing	solver=yes	solver-maybe	solver=yes	solver=yes	8	these	others

[remove following 0 transformations](#)

Figure 6.3: screenshot of the experimental query interface

7

Implementation

The Star-Exec-Presenter is a web application written in Haskell and based upon the Yesod Web Framework. In this chapter I will talk about the actual implementation of the tool as well as about the benefits of using especially Haskell and Yesod.

7.1 Haskell

“Haskell is a general purpose, purely functional programming language incorporating many recent innovations in programming language design. Haskell provides higher-order functions, non-strict semantics, static polymorphic typing, user-defined algebraic datatypes, pattern-matching, list comprehensions, a module system, a monadic I/O system, and a rich set of primitive datatypes, including lists, arrays, arbitrary and fixed precision integers, and floating-point numbers. Haskell is both the culmination and solidification of many years of research on non-strict functional languages.” (Peyton Jones 2003, Marlow (2010))

This quote stems from Simon Peyton Jones and his paper “Haskell 98 language and libraries: the revised report” as well as from the paper “Haskell 2010 Language Report” by Simon Marlow. It does sum up the important features of the Haskell programming language. In addition I want to add that Haskell has type inference and its type system is very expressive. In the following I try to explain the listed terms.

Haskell is a general purpose programming language which means it can be used to

develop a web-server, a desktop application as well as a simple commandline tool. It is suited for a wide range of application domains. Haskell is a purely functional programming language, i.e. it follows the functional programming paradigm and every pure function has no side effect. Functional programming relies on a mathematical approach where every calculation is defined by expressions. Everything is an expression, whether it's a function or a value. Purely means, that a function only works with its input arguments and returns the same result no matter how often it is called with the same arguments. Every expression is immutable so once they are defined they cannot be changed anymore. Haskell code therefore is easily testable and even proofable. In contrast, imperative programming languages allow mutability and functions are subroutines that can have side effects. A side effect for instance can be a simple tracing that doesn't effect the function's result or it can be the change of a global state.

```
1 -- example of a pure function
2 mul a b = a * b
3
4 -- example of an impure function
5 mulM a b = do
6     putStrLn ("multiplying " ++ (show a) ++ " with " ++ (show b))
7     return (a * b)
```

It is important to note that, unlike programming languages like Java or C, Haskell is indentation based. So, there are no blocks surrounded by curly brackets, but every whitespace at the beginning of each line has a meaning to the program.

7.1.1 Higher-Order Functions

Haskell has higher-order function and emphasizes their usage. Higher-order functions are functions that require functions as input parameters or return a new function. Therefore functions are data that can be passed round. Thanks to currying, all functions are higher-order functions and thereby expressions. Currying means that a function always returns a new function if it isn't called with the full list of arguments.

```

1 -- example for a higher-order function
2 map :: (a -> b) -> [a] -> [b]
3 map _ []      = []           -- special case for an empty list
4 map f (x:xs) = f x : map f xs -- f applied to the list's head
5                                     -- recursive call with the list's
                                     rest

```

7.1.2 Pattern-Matching

The example from above also shows pattern-matching, which is a rich programming technique to identify certain cases for a given expression. The `map` function from above has two implementations, one for the case of the empty list which also is the base case for the recursion, and the main implementation that declares the algorithm of the function. Pattern matching helps simplifying code by reducing the usage of `if`-statements and alike.

7.1.3 List Comprehensions

To stay with the `map` function, list comprehensions are another way to define this method as well as lists in general. They are similar to set comprehensions in mathematics. For example the list comprehension `[2*x | x <- [1..10]]` is equivalent to the set comprehension $\{2x | x \in \mathbb{N}, x \leq 10\}$. (Lipovaca 2012) Both expressions result in a list of even natural numbers from two to 20. Referencing the `map` function from above, an alternative implementation follows:

```

1 map f xs = [ f x | x <- xs ]

```

Another example of list comprehensions is the solution to the first problem of Project Euler¹. The exercise is to find the sum of all multiples of three or five below 1000:

```

1 sum [x | x <- [3..999], x `mod` 3 == 0 || x `mod` 5 == 0]

```

7.1.4 Non-Strict Semantics, Lazy Evaluation

Another essential feature of Haskell is the non-strict semantics. Non-strict in terms of Haskell means that every expression is evaluated by need. So, only declaring

¹see <https://projecteuler.net/problem=1>

an expression doesn't invoke its evaluation until it is needed, e.g. for output. For instance the fibonacci sequence can be seen as an infinite list. This sequence can also be implemented via a list comprehension:

```
1 let fibs = 0 : 1 : [ a + b | (a, b) <- zip fibs (tail fibs)]
```

Here, we have the declaration of an infinite list, but which will be evaluated only if it is needed. These non-strict semantics in Haskell are implemented as *lazy evaluation* which leads to the expression from above also being evaluated by need. So until it is really needed, `fibs` is only a thunked² expression. But this strategy can lead to high memory usage especially for complex algorithms, so in some cases it is advised to use strict evaluation which can be used with the `seq` function in Haskell's `Prelude` module. Other modules with strict evaluations for example are `Data.Map.Strict` or `Data.List`, the latter having a strict implementation of the `Prelude`'s `foldl` function. Especially this function can lead to a stack overflow with a large list.³

7.1.5 Type System, Type Inference

Haskell has a rich type system which can automatically infer the type of an expression. For example the fibonacci sequence from above has the type `fibs :: Num b => [b]` which was inferred by the `(+)` operator's type `((+) :: Num a => a -> a -> a)`. Not all types can be inferred and the compiler outputs a compiler error, so a type signature should be added. In most cases these signatures are optional but also can help improving the readability of the code. These signatures don't contain any special type information because they are polymorphic, that is, they can be used with any type that meets the functions requirements, may it be `Int` or `Double`. Polymorphic types are denoted by small alphanumeric characters like `a`, `b` or `t0`. They can be restricted by a type-class which is a set of functions. The `(+)` operator is a function of the `Num` type-class which is a set of functions to operate on numerical values. Type-classes are similar to interfaces of imperative programming languages like Java.

7.1.6 Algebraic Datatypes

With Haskell a developer can declare custom algebraic datatypes. What that means is that a type can be defined by the algebraic operations *sum* and *product*, where a *sum*

²a value that is yet to be evaluated

³see https://www.haskell.org/haskellwiki/Foldr_Foldl_Foldl%27

is a group of alternative data constructors and *product* is the combination of them. Below are some examples:

```
1 -- sum of polymorphic types
2 data Either a b = Left a | Right b
3
4 -- product of polymorphic types
5 data Pair a b = Pair a b
6
7 -- a record datatype
8 data Record a b =
9     Record
10     { memberA :: a
11       , memberB :: b
12     }
```

The `Either` datatype is a special type that is used for computations that can result in an error, where the `Left` constructor contains the error and the `Right` constructor holds the successful result. This leads to Haskell's expressiveness which is one of the benefits of its type system. The `Either` datatype explains that a function can return an error, so a developer has to consider the returned value. Java in contrast is more verbose by forcing to add a `throws` statement to the method definition. A developer has to use a `try ... catch` block to manage an error. Another example for the expressive type system of Haskell is the `Maybe` datatype whose signature follows:

```
1 data Maybe a = Nothing | Just a
```

This datatype is used for computations that can have no result or that require optional arguments. The `Nothing` constructor means that there is no value, whereas the `Just` constructor wraps a result or an optional argument. The `Maybe` datatype is useful because it prevents undefined or null values via the type system. In Java, for example, carelessly developed programs can pass the compiler but also result in a `NullPointerException` whose cause can be difficult to resolve. Of course, that doesn't mean that a Haskell program doesn't have errors, but the type system and the compiler helps to reduce them or even prevent them all.

7.1.7 Modules

Haskell has a distinct module system, where each module contains a well-matched set of functions, types and type-classes. The Haskell code base is separated into a large number of modules, each one serving a certain kind of purpose. Each Haskell program has a main module which contains the main function as well as a set of other modules. It is important to know that, without any additional language extensions, a bunch of modules cannot be defined circularly, that is module A cannot import module B if this already imports module A.

7.1.8 Input/Output, Monads

All these features and benefits of Haskell are of no use if it can't communicate with the real world, that is doing input/output. So, Haskell has an appropriate I/O system which is also monadic⁴. That means that every function that operates within the IO-Monad has a special type signature. Below are some examples:

```
1 putChar      :: Char -> IO ()
2 putStrLn    :: String -> IO ()
3 getLine     :: IO String
4 readFile     :: FilePath -> IO String
5 writeFile   :: FilePath -> String -> IO ()
```

All functions return a value within the context of the IO-Monad. To access the content it must be bound to a function that processes it. For example, the `putStrLn` function can be implemented via the `putChar` function and a monadic composition (`(>>)`):

```
1 putStrLn' s = mapM_ putChar s >> putChar '\n'
```

This implementation uses the monadic version of the `map` function to put each `Char` of the given `String` to the output followed by a newline. To simplify the work with the IO-Monad or Monads in general Haskell has the `do`-notation which I want to explain with the following example:

⁴A Monad is an abstract datatype, that wraps a (primitive) value in a context and binds it to a function operating within that context


```

1 func_do f g h = do
2     a <- f           -- binds the wrapped value of f to a
3     b <- g           -- binds the wrapped value of g to b
4     c <- h           -- binds the wrapped value of h to c
5     return (a, b, c) -- wraps the result in the default context
6
7 func_bind f g h =
8     f >>= \a ->
9         g >>= \b ->
10             h >>= \c ->
11                 return (a, b, c)

```

Both functions essentially do the same. They take three monadic values (e.g. `Maybe a`) and bind the content of these values to `a`, `b` and `c`. The first function uses `do`-notation, the second one uses the `bind`-Operator (`(>>=)`). The `do`-notation is a sugared way of using the `bind`-Operator, i.e. the compiler transforms the code of `func_do` to the code of `func_bind`. It is the same with all functions of the `IO-Monad`. But there is one important note: The `IO-Monad` is impure. What does this mean? All functions in Haskell are pure unless they use the `IO-Monad`. Because this `Monad` communicates with the real world, Haskell has no control of this communication. That means, that every call of an `IO-function` the result is not guaranteed to be the same. A File could have changed since the last time it was read. Another good example is the operation system's random number generator. Each time, it produces another value, so this process can never be pure. (O'Sullivan, Goerzen, and Stewart 2010)

7.2 Yesod Web Framework

The Yesod Web Framework is a set of tools and libraries build with Haskell. It aims to improve the development of server-side, RESTful web applications by providing type-safety, conciseness and performance. It is divided into several components that together form the final web application. Yesod consists of an object-relational mapping for the communication with a database, several domain-specific languages (DSLs) for routes, templates and the database types, it has session- and forms-handling as well as the support for authentication, authorization or internationalization. All DSLs are implemented with Template Haskell, an extension to the

GHC⁵ to enable metaprogramming. Template Haskell code will be transformed to actual Haskell code while compiling. (Sheard and Jones 2002)

The basis of Yesod is the Warp server which is a fast web server build in Haskell. It is the main implementation of the Web Application Interface (WAI) which in turn is a generalized interface for building web server. It generally considers the nature of the web infrastructure, where the communication essentially works like a function: a client sends a request and gets a response in return. This communication can be illustrated with the following type signature:

```
1 askServer :: Request -> Response
```

WAI aims to be an efficient interface by relying on lazy IO, which leads to a small memory footprint. An example of a small WAI implementation follows below:

```
1 {-# LANGUAGE OverloadedStrings #-}
2 import Network.Wai
3 import Network.HTTP.Types (status200)
4 import Network.Wai.Handler.Warp (run)
5
6 application _ respond = respond $
7   responseLBS status200 [("Content-Type", "text/plain")] "Hello
8     World"
9 main = run 3000 application
```

This program takes a request and responds with the text “Hello World”, regardless of what the request may contain. A more complex example considers the requested path (route) and responds according to that:

⁵Glaskow Haskell Compiler, the most used compiler infrastructure for Haskell

```

1 {-# LANGUAGE OverloadedStrings #-}
2
3 import Network.Wai
4 import Network.Wai.Handler.Warp
5 import Network.HTTP.Types (status200)
6 import Blaze.ByteString.Builder (copyByteString)
7 import qualified Data.ByteString.UTF8 as BU
8 import Data.Monoid
9
10 main = do
11   putStrLn $ "Listening on port 3000"
12   run 3000 app
13
14 app req respond = do
15   let res = case pathInfo req of
16     ["yay"] -> yay
17     ["foo", "bar"] -> foobar
18     x -> index x
19   respond res
20
21 yay = responseBuilder status200
22   [ ("Content-Type", "text/plain") ]
23   $ mconcat $ map copyByteString [ "yay" ]
24
25 index x = responseBuilder status200
26   [("Content-Type", "text/html")] $ mconcat $ map copyByteString
27   [ "<p>Hello from ", BU.fromString $ show x, "!</p>"
28     , "<p><a href='/yay'>yay</a></p>\n" ]
29
30 foobar = responseBuilder status200
31   [("Content-Type", "text/html")] $ mconcat $ map copyByteString
32   [ "<p>Hello from foobar! ;)</p>" ]

```

7.2.1 Routes

Yesod has a more complex mechanism for handling requests build atop of WAI and Warp. The routes are defined via a special DSL which concentrates on the path and method of each request. A path is divided into one or more path pieces each one being static or dynamic, where dynamic parts are represented by a distinct type. There can be a single dynamic path piece as well as any number. Each path is followed by the handler resource and the supported method which can be a standard HTTP method or even a custom one. The handler resource is the connection to the Haskell code. Template Haskell generates a handler function for each route and method as well as special datatype representing the route. This datatype enables typesafe URLs. In the section *6.3 REST interface* I showed an example of the routes DSL. Below is an explaining example:

```
1 -- the home route
2 -- generates the handler "getHome"
3 -- "HomeR" is the datatype for this route
4 /                HomeR      GET
5
6 -- the list of users
7 -- generates the handler "getListUsers" and "postListUsers"
8 -- "ListUsersR" is the datatype for this route
9 /users           ListUsersR  GET POST
10
11 -- generates the handler "getUser" and "postUser"
12 -- "UserR UserId" is the datatype for this route
13 /users/#UserId  UserR       GET POST
14
15 -- generates the handler "getList"
16 -- "ListR [Text]" is the datatype for this route
17 -- "Texts" is a type synonym for "[Text]"
18 /list/*Texts    ListR       GET
19
20 -- special routes for static files and authentication
21 /static StaticR Static getStatic
22 /auth   AuthR   Auth   getAuth
```

7.2.2 Templates

Yesod uses templates for HTML, JavaScript and CSS⁶ as type-safe expressions which are inserted into the application at compile time. Again, Template Haskell takes care of them. Templates are inserted into the handler functions via Quasi-Quoters, some of which take the template code while others need to get the filename of the template. So, Yesod is very flexible.

In Yesod, all template DSLs are named after a Shakespear character. For HTML, there is the Hamlet markup language:

```
1 <div>
2   <h1>#{myTitle}
3   <p>
4     <a href="@{MyCustomRouteR}">my custom route</a>
5   $maybe val <- maybeVal
6   <p>val
7   $nothing
8   <p>There is Nothing to show here
9   $if null someList
10  <ul>
11    $forall item <- someList
12    <li>#{show item}
13  $else
14    <p>The List is empty
15 ~{footerWidget}
```

As you can see, the syntax is very similar to real HTML. But, it is important to note, that Hamlet is mostly indentation based. *Mostly* means, that indentation is used for block elements like `div`, `p` or headlines, so these elements doesn't have a closing tag in Hamlet. Actually they would cause a compiletime error. Inline elements such as anchor tags (``) must have a closing tag. The integration with actual Haskell code has several incarnations. The first example is the interpolation of the expression `myTitle` which is bracketed by `#{...}`. It will be inserted into the the HTML code with the `toHtml` function of the `ToHtml` type-class. So every expression must be an instance of this type-class to be interpolated. The next feature of Hamlet is typesafe URLs which incorporates with the routes definition. The resource datatypes from the routes DSL can be used here. So an expression bracketed by

⁶Cascading Stylesheet

`@{...}` generates a typesafe URL within a `href` attribute. This is possible for all attributes containing a path to a resource.

Another feature is the special treatment of `Maybe` values as well as lists. Here, special statements initialize Haskell code which handles the values. So a `Maybe` value will be unwrapped (`$maybe`) or an optional information can be shown if there is `Nothing` (`$nothing`). The same applies to lists where the statement `$forall` indicates a loop to process all items of a list. Additionally, there are logical statements like `$if/$else` as well as `$case/$of`, the latter also supporting pattern matching. And to shorten inconveniently long expressions a new one can be declared via `$with shortExp <- longExp`.

The shakespearean templates are round out with the DLSs Julius for JavaScript and Lucius as well as Cassius for CSS. Julius is basically JavaScript with the features of interpolation. Lucius and Cassius are both an optimized subset of CSS equivalent to the likes of SASS⁷ or Less⁸. They feature nested blocks, interpolated expressions, variable declaration as well as mixins. Mixins are reusable code blocks, e.g. for dealing with vendor prefixes. Whereas Lucius is more like original CSS with curly brackets enclosing the CSS statements, Cassius is indentation based. Below is a code example for Lucius taken from Star-Exec-Presenter:

```
1 @colorYes: #80FFB0;
2 ...
3
4 .table > tfoot > tr {
5     > td.solver-yes,
6     > th.solver-yes {
7         background-color: #{colorYes};
8     }
9     ...
10 }
```

7.2.3 Persistence

For the communication with and the integration of a database, Yesod relies on the persistent framework. This framework enables a generalized way of working with the database regardless of its kind, may it be SQL or NoSQL. Even for persistent

⁷a dynamic stylesheet language for CSS (<http://sass-lang.com/>)

⁸a dynamic stylesheet language for CSS (<http://lesscss.org/>)

datatypes there is a DSL which simplifies their declaration as there are many special types incorporating. Below is an example:

```
1 User
2   ident Text
3   password Text Maybe
4   name Text
5   gender Gender
6   UniqueUser ident
7   deriving Show
```

This is a simple example of representation of a user in the database. Just like most DSLs, this one is indentation based, too. It declares the type `User` with four members and a uniqueness constraint. Each member is defined by its name followed by its type. Unlike Haskell where an optional type is defined by `Maybe a`, the persistent DSL requires `Maybe` to be at the end of the declaration. A unique constraint is defined with the prefix `Unique` to the type's name as well as the names of one or more members of the type. A full description of the persistent DSL syntax can be found in the wiki of the `persistent`-repository on GitHub⁹.

From this definition several types are generated. First and foremost a record datatype will be defined representing the actual type. This datatype is made instance of several type-classes to make them compatible with the persistent interface. Additionally a data constructor is available to represent a unique instance of the data in the database. There are also several datatypes for filtering the data table by specific constraints. Custom types, such as `Gender` must be made an instance of the type-class `PersistField` which takes care of serialization and deserialization of the custom datatype. An example outcome follows:

```
1 data User = User { userIdent :: Text
2                   , userPassword :: Maybe Text
3                   , userName :: Text
4                   , userGender :: Gender }
5   deriving (Show)
6
7 -- datatype representing the primary key of the table
8 -- will be serialized to the actual database type value
9 data UserId = KeyBackend SqlBackend User
```

⁹see <https://github.com/yesodweb/persistent/wiki/Persistent-entity-syntax>

To actually communicate with the database the type-classes `PersistStore`, `PersistUnique` and `PersistQuery` are used. All three offer useful functions to query and store all related data in the database. The following listing shows some possible usages:

```
1 import qualified Data.Text as T
2
3 communicateWithDB = runDB $ do
4     let ident = T.pack "haskell@example.org"
5         name = T.pack "Haskell Curry"
6         gender = Male
7         user = User ident Nothing name gender
8     -- inserting the user entity
9     -- but could cause an error if the user's already in the db
10    userId <- insert user
11    -- because of uniqueness constrain
12    -- insertUnique is better suited
13    mUserId <- insertUnique user
14    case mUserId of
15        Just userId -> print userId
16        Nothing      -> error "user already in db"
17    -- setting the password of the user
18    let passwd = Just $ T.pack "someSecretPassPhrase"
19    update userId [ UserPassword =. passwd ]
20    -- get unique user entity
21    mUser <- getBy $ UniqueUser ident
22    case mUser of
23        Just user -> print user
24        Nothing   -> error "No user found"
25    -- counts all users by a certain filter
26    numberOfMales <- count [ UserGender ==. Male ]
27    -- fetches 10 users by a certain filter
28    femaleUsers <- selectList [ UserGender ==. Female ]
29                        [ Asc UserName
30                        , LimitTo 10
31                        , OffsetBy 0 ]
32    -- deleting a complete table
33    deleteWhere ([] :: [Filter User])
```


The `runDB` function wraps all actions within a transaction. Actually, each call of `insert` or `selectList` can be called on its own with `runDB`. But to make the code more efficient, it is useful to put all steps into one call to `runDB`. Another advantage is that an error rolls back the full transaction.

(Snoyman 2012)

8

Evaluation

In this chapter I want to evaluate the previously set goals. How did Star-Exec-Presenter perform during the Termination Competition 2014 and which problems have to be discussed in the future. Also, are there any considerations to be made about StarExec?

8.1 Termination Competition 2014

First of all, the 2014 Termination Competition can be seen as a success. The competition run passed without any failures and Star-Exec-Presenter mostly if not always ran stable during the event. We also got positive comments from the community.¹² The tool was able to handle multiple requests to StarExec as well as processed the data according to the requirements of the competition. Star-Exec-Presenter proved to be a huge first move towards being a useful tool for the Termination Community. Nearly all requirements were met in their basic sense.

To give some numbers, the database for our tool allocates about 140 Megabytes. The code base of Star-Exec-Presenter consists of about 45.000 lines of code. The compilation took about 25 Minutes in a virtual machine utilizing two CPU-Cores

¹see <http://lists.lri.fr/pipermail/termtools/2014-July/000971.html>

²see <http://lists.lri.fr/pipermail/termtools/2014-July/000972.html>

and four Gigabytes main memory on my MacBook Pro (2012) with a 2.9 GHz Core i7 Processor.

Of course, there is much room for improvements. The RESTful interface can be even more RESTful, e.g. that it also can respond with an XML or JSON output. The web interface is also very much basic and not really functional. In an optical sense Star-Exec-Presenter could have more of an own identity. And most importantly, as it still is experimental, the querying could be more elaborated. The web interface isn't very useful and the query-URLs are only serializations/deserializations of the underlying datatypes.

8.2 Future considerations

For future versions of Star-Exec-Presenter, there are some points I want to propose in some ideas. Currently, the caching mechanism is invoked by a client's request. Additionally a background worker can handle the whole caching alone by regularly traversing the database's content and requesting the latest information from StarExec. Also, the internal caching stores can be defined better. At the moment, the database is not only used for the data from StarExec but also for meta information such as which information is currently requested. This information should be better put in the main memory via STM rather than in the database.

Of course, the whole code quality of Star-Exec-Presenter could be improved and the expressiveness of Haskell could also be better used throughout the project. This leads to a major improvement: a more general basic code base to be used not only by the Termination Community. But, it may be difficult to meet all the requirements of all the other communities.

Another aspect is the web interface itself, whose usability could probably be improved by the utilization of a frontend framework such as AngularJS³ or Ember.js⁴. With the help of AJAX the web interface could feel more like an application as it is possible to completely avoid any page reloads. A more important issue is the export of all the data stored in the database. So, a generalized export could be useful as we currently have three types of data sources: StarExec, LRI and UIBK. If a new instance of Star-Exec-Presenter will be installed, all these exported data could then be imported again. Of the same importance is the import which in one aspect

³see <https://angularjs.org/>

⁴see <http://emberjs.com/>

isn't finally implemented. This aspect regards the identification of the benchmarks throughout all data, whether its directly from StarExec or from imported sources. Currently there is no mapping whatsoever, but it is a key to comparing job results from different competitions.

8.3 StarExec

The StarExec service is a crucial part of the Star-Exec-Presenter. Without it, there would only be the old data for use. During the work on our tool, we experienced some problems, which I want to point out. Hopefully, this may lead to improvements by – or at least inspirations for – the developers of StarExec. First of all, there were times when StarExec was inaccessible. Luckily, this happened only during the beginning of the development of Star-Exec-Presenter. So, this might already be improved. Another area that could profit from improvements is the whole interface by which we collect all the data. Currently, it is a mix of zipped CSV files, HTML webpages or some JSON. There could be a RESTful API that is designed for tools such as Star-Exec-Presenter. In the following I want to propose such an API (which is loosely based on our interface):

1 /jobs/#JobID	GET
2 /solvers/#SolverID	GET
3 /benchmarks/#BenchmarkID	GET
4 /pairs/#JobPairID	GET
5 /post_procs/#PostProcID	GET
6 /results/#JobID	GET

By supporting JSON, CSV or even XML and HTML, this API could meet any needs. Of course, this has to be implemented by the developers of StarExec who might not have enough resources available to fulfill our wishes.

9

Summary

The goal of this thesis was to give a brief introduction to the Star-Exec-Presenter, a web-based application for data-mining results from StarExec. At first, I gave an introduction to the *Termination Community* and its subject, the *Termination of Term Rewriting*. I explained *StarExec* and its functionalities in terms of our tool and listed the *Use Cases* which we wanted to be more accessible. The *Requirements* defined the foundation of the *Design* of Star-Exec-Presenter and its *Implementation*. In the *Evaluation* I tried to be self-critical in what I achieved with this whole project and gave *Future Consideration* how our tool can evolve in preparation of the next Termination Competition or a possible utilization by another community.

Glossary

AJAX

Asynchronous JavaScript and XML, a technique to update a web-page without fully reloading it by fetching only the needed data.

Algebraic Datatype

A datatype created by the algebraic operations of *sum* and *product*.¹

API

Application Programming Interface, a set of components (methods, protocols, tools) for using a software.

Benchmark

A certain problem to be processed by a solver.

Caching

A fast store to put data in, that can then be requested much faster.

CSS

Cascading StyleSheet, a textbased format to describe the appearance of HTML.

CSV

Comma Separated Values, a textbased format that describes a table mit its content. Each column is seperated by a comma or an otherwise specified seperator.

DSL

Domain Specific Language, a programming language to serve a certain purpose. A DSL is often less powerfull than a general purpose programming language.

¹see https://www.haskell.org/haskellwiki/Algebraic_data_type

GUI

Graphical User Interface, a special type of interface which a user can actually see and use.

Haskell

A purely functional programming language with an expressive type system that supports type inference, pattern matching and lazy evaluation.

Higher-Order Function

A function that can be treated like data und therefore can be passed round.

HTML

HyperText Markup Language, a structural textbased language to describe content which will be interpreted and displayed by a web-browser.

HTTP

HyperText Transfer Protocoll, a standard way to communicate with the aid of a web browser.

JavaScript

A dynamic programming language originally created for dynamic content in web browsers. It serves imperative as well as functional programming paradigms.

Job

A complex task on StarExec, that starts solvers and stores their results.

Job-Pair

A unique combination of a certain benchmark and a certain solver within a job. It contains the actual output of a solver according to the processed benchmark.

JSON

JavaScript Object Notation, a lightweight data-interchange format originally specified by Douglas Crockford. It is a subset of JavaScript as well as a human-readable data-format.²

Lazy Evaluation

A non-strict evaluation strategy, expecially used by Haskell. An expression won't be evaluated until it is really needed, e.g. for output.³

²see <http://www.json.org/>

³see https://www.haskell.org/haskellwiki/Lazy_evaluation

Monad

An abstract datatype, that wraps a (primitive) value in a context and binds it to a function operating within that context.

MVC

Model View Controller, a software design pattern where data, logic and presentation are separated from each other.

Pattern Matching

A rich programming technique to identify certain cases for a given expression. It helps simplifying code by reducing the usage of if-statements and alike.⁴

REST

Representational state transfer, a paradigm to implement a server-client communication.(Fielding 2000)

StarExec

A cross community logic solving service developed at the University of Iowa.⁵

Star-Exec-Presenter

A Haskell web application to data mining, displaying and querying competition results from StarExec.

SOAP

Simple Object Access Protocol, a network protocol for data-interchange between remote systems.

Solver

A tool that processes a benchmark on StarExec.

SQL

A relational database query language.

STM

Software Transactional Memory, a mechanism for concurrent programming with shared resources.

⁴see <https://www.haskell.org/tutorial/patterns.html>

⁵see <https://www.starexec.org/starexec/public/about.jsp>

Termination

The final state of a program, with which a program ends if it ends.

Termination Community

An international group of developers and researchers working in the field of Termination of Term Rewriting.⁶

Term Rewriting System

An alternative computational model to a Turing machine.(Zantema 2000)

TPDB

Termination Problem Data Base, a collection of problems regarding the Termination of Term Rewriting.

Type Inference

Automatic detection of a type of a certain expression.

URI

Uniform Resource Identifier, an identifier for certain resources especially for the world wide web.

URL

Uniform Resource Locator, a subset of URIs used for web-sites.

WAI

Web Application Interface, a general interface to write web servers with Haskell.⁷

Warp

A fast web-server used by the Yesod Web Framework programmed in Haskell.

Yesod Web Framework

A web framework programmed in Haskell and initially founded by Micheal Snoyman.(Snoyman 2012)

⁶see http://www.termination-portal.org/wiki/Main_Page

⁷see <http://www.yesodweb.com/book/web-application-interface>

References

- Fielding, Roy Thomas. 2000. *Architectural Styles and the Design of Network-Based Software Architectures*. <http://www.ics.uci.edu/~fielding/pubs/dissertation/top.htm>.
- Lipovaca, Miran. 2012. *Learn You a Haskell for Great Good!: a Beginner's Guide*. San Francisco, CA: No Starch Press.
- Marché, Claude, and Hans Zantema. 2007. “The Termination Competition.” In *In Proc. RTA '07, LNCS 4533*, 303–13.
- Marlow, Simon. 2010. “Haskell 2010 Language Report.” <https://www.haskell.org/definition/haskell2010.pdf>.
- . 2013. *Parallel and Concurrent Programming in Haskell*. First edition. Beijing: O'Reilly.
- O'Sullivan, Bryan, John Goerzen, and Donald Bruce Stewart. 2010. *Real World Haskell*. Auflage: 1. Sebastopol, CA: O'Reilly.
- Peyton Jones, Simon L., ed. 2003. *Haskell 98 Language and Libraries: the Revised Report*. Cambridge, U.K. ; New York: Cambridge University Press.
- Richardson, Leonard. 2007. *RESTful Web Services*. Farnham: O'Reilly.
- Sheard, Tim, and Simon Peyton Jones. 2002. “Template Meta-Programming for Haskell.” In *In Proceedings of the ACM SIGPLAN Workshop on Haskell*, 1–16. ACM.
- Snoyman, Michael. 2012. *Developing Web Applications with Haskell and Yesod*. 1st ed. Sebastopol, CA: O'Reilly.
- Stump, Aaron, Geoff Sutcliffe, and Cesare Tinelli. 2014. “StarExec: A Cross-Community Infrastructure for Logic Solving.” In *Automated Reasoning*, edited by Stéphane Demri, Deepak Kapur, and Christoph Weidenbach, 8562:367–73. Lecture Notes in Computer Science. Springer International Publishing. http://dx.doi.org/10.1007/978-3-319-08587-6_28.

Zantema, Hans. 2000. *Termination of Term Rewriting*.

Eidesstattliche Versicherung

Ich erkläre hiermit, dass ich diese Masterarbeit selbstständig ohne Hilfe Dritter und ohne Benutzung anderer als der angegebenen Quellen und Hilfsmittel verfasst habe. Alle den benutzten Quellen wörtlich oder sinngemäß entnommenen Stellen sind als solche einzeln kenntlich gemacht.

Diese Arbeit ist bislang keiner anderen Prüfungsbehörde vorgelegt und auch nicht veröffentlicht worden.

Ich bin mir bewusst, dass eine falsche Erklärung rechtliche Folgen haben wird.

Leipzig, der 23.01.2015