Master Thesis

A Haskell Web-Application for Data-Mining Competition-Results from StarExec

A thesis submitted to attain the degree of
Master of Science at the Leipzig University of Applied Sciences
(M.Sc. HTWK Leipzig)

presented by
Stefan von der Krone

B.Eng. in Media Technology, University of applied Sciences Mittweida

born on 19th of April, 1984

citizen of Germany

accepted on the recommendation of Prof. Dr. Johannes Waldmann

2014

Table of Contents

1
Introduction

This is a guide to the Star-Exec-Presenter, a web-based visualization and data-mining software developed in preparation of this thesis. The Star-Exec-Presenter is intended to be an efficient to use software for members of the termination community of the StarExec service, as well as other communities or users of StarExec. The web application features essential functionalities such as starting competitions, loading results and other data from StarExec or visualizing and filtering such data.

Termination is a branch of theoretical computer science where programs such as algorithms are examined, whether they terminate, that is whether they complete. In terms of StarExec, solvers are running over a bunch of problems returning an answer for each problem: Yes, No or Maybe.

The Termination Competition is an annual event of the termination community of StarExec, where solvers from different contributors compete in several categories. In 2014 the competition was organized by my adviser, Prof. Johannes Waldman, and me. Our goal was to do both, starting the competition run as well as to present it via an automatically updated web-interface.

The 2014 competition caused a total amount of 13 Gigabyte of data and needed about 90 days of CPU time on all of the 192 Quad-Core CPUs of the StarExec cluster. (Waldmann 2014)

Our additional goal was to be able to compare these results with those from previous

ones. So we added an option to import the data from the 2007 competition hosted by the Laboratory of Computer Science at Université Paris-Sud and the following competitions hosted by the University of Innsbruck.

"StarExec is a cross community logic solving service developed at the University of Iowa (…)." (StarExec 2013) It is a technical infrastructure providing the service to run logic solvers on a powerful cluster of CPUs. It also provides an extensive web-based user interface to upload and run solvers and problems, referred to as benchmarks. A REST-based[1] API[2] is utalized by the Star-Exec-Presenter.

The Star-Exec-Presenter is the software developed as the concrete work for this thesis. It is a web application programmed in Haskel[3] with the utalization of the Yesod Web Framework[4].

The motivation of this thesis is to provide a tool for StarExec users for further research as well as some kind of standardization of hosting and running future termination competition and competitions of other StarExec communities. The goal for Star-Exec-Presenter is to be a tool, that is easy to use as well as easy to install and run. It is open source, so it can be forked und changed to better meet future needs.

With Star-Exec-Presenter we wanted to extend the ability of the StarExec service with filtering of the results and comparing them with older ones from previous competitions. Contributers of solvers should be able to track their progress of development of their respective solver through time.

---

[1]Representational state transfer, a paradigm to implement a server-client communication

[2]Application programming interface, a set of components (methods, protocols, tools) for using a software

[3]a functional, non-strict, declarative programming language (http://www.haskell.org/)

[4]a REST-based web application framework (http://www.yesodweb.com/)

2

Termination Competition

In this chapter I will discuss the terms of Termination and Termination of Term Rewriting. I also will describe the Termination Competition, as its 2014 incarnation was one of the goals towards the Star-Exec-Presenter was developed.

Termination of Term Rewriting

(Zantema 2000)

Termination Community

Termination Competition

(Marché and Zantema 2007)

3

StarExec

StarExec is the service that is utilized by the Termination Community and other communities to run their research as well as the latest Termination Competition. This chapter will examine the interfaces from which the relevant data is fetched as well as their special characteristics. I will go into the details of how StarExec works and how to get the data.

Interfaces

4

Use Cases

In this chapter I will list the use cases which the Star-Exec-Presenter is mapped onto. These use cases range from simple tasks, for instance displaying solver-results from StarExec, up to more complex ones, like starting and monitoring a competition.

Displaying Results and Outputs

The core of StarExec is not only to run the solver on specific benchmarks but also to host the results and outputs of each run. So a major task for the Star-Exec-Presenter is to display that data in an appropriate way. The explicit data is the result of a solver with a specific benchmark and its CPU- as well as Wallclock-Time. Additionally the solver's output is important which is accessible on StarExec with the job-pair, a unique combination within the specific job of solver and benchmark.

Compare Results

Simply displaying the results and outputs of each solver-run is not enough. It is more important to compare the results between several solver. Are there differences between their results on specific benchmarks? Also, a list of jobs should be displayable

at once. A historical comparison is the most important aspect of this use case, as the developers of a solver would want to track the evolution of their tool.

Querriyng

As important comparing the result is as overwhelming is their amount, so a beneficial use case is querriyng the results according to specific filters. Such filters can eliminate unimportant results from the representation. For instance a benchmark is uninteresting if all solvers have the same result, so all those benchmarks can be filtered out.

The user should be able to concentrate on data which meets his or her interests. This data could tell which benchmark causes incorrect results with the own solver. Also it can illustrate changes in the solvers results in comparison to past runs, fo instance a once solved benchmark could now be unsolvable. This is the data we want to querry in terms of data mining.

Interpret Outputs as possible XML-Proofs

Some solver explicitly log an XML-proof with each run in its output. This proof should be detected as well as extracted from the output. Also the respective XSL-File should be used to display the proof in the browser.

Start a predefined Competition

One of the core use cases of the Star-Exec-Presenter is the ability to start a competition run on StarExec. As StarExec actually doesn't know what a competition in our terms is, Star-Exec-Presenter should know it. So all meta-information of a competition should be managed by the Star-Exec-Presenter. These meta information include hierarchical information of the competition's organisation as well as the details which we'll send to or get from StarExec.

Hierarchical information is meant as how each solver and benchmark is organised within the competition. Each competition consists of several meta categories which cover a general topic in the competition. The 2014 Termination Competition has the

following: "Termination of Term Rewriting (and Transition Systems)", "Complexity Analysis of Term Rewriting" and "Termination of Programming Languages".

Each meta category is devided into smaller, more specific categories, each being a representation of an actual job on StarExec. These categories may be best described by the meta categories "Termination of Programming Languages", as there is a category about the programming language C. Other categories could bring up Java, Logic Programming or Functional Programming.

Displaying a Competition and its Results

The Star-Exec-Presenter is a tool that ran during the 2014 Termination Competition. It should display the current status of the competition run with a compact web interface that automatically updates itself.

5

Requirements

This chapter discusses the requirements that were agreed to. Which data is interesting and how is it handled? How should the application be build, which kind of data should be accessible through the Star-Exec-Presenter? And how should the application incorporate with StarExec? All requirements are infered from the previous chapter.

Solver and Post processor output

To display the results of a job on StarExec we must now which data actually is important. So, basically StarExec knows two different kinds of resulting data. The first is a table which contains the resulting data of each solver-benchmark combination (the job-pair) of the job. It tells us the final result of a job-pair as well as its CPU- or Wallclock-Time and its current status. Each job-pair has a unique identifier which leads us to the second kind of resulting data: the output of a job-pair which is accessible separately. It contains the standard-out of the solver and a log generated by StarExec.

The aforesaid table should be displayed wellarranged with the Star-Exec-Presenter. Important data kinds are of course the result of the job-pair as well as the CPU- and Wallclock-Time. The presentation should be easy to read and its information

gathered quickly.

The output of the job-pair should be scanned for XML to determine wether it contains a proof. The corresponding presentation should give the chance to display that proof according to its required XSL-Stylesheet.

All data should be persistently stored in a Postgre-SQL database.

Compact web interface

The web interface of StarExec is quite verbose and not specifically matching the needs of members of the termination community. So, the Star-Exec-Presenter should feature a much more compact web interface which clearly guides the user to the destined information. Readability is important.

Short response time

As StarExec is located in Iowa and some requests require expensive fetching from its database some responses may need up the seconds to be fully available. This is not desired as some users may consider it an error of the Star-Exec-Presenter. So our tool should cache all data that it loads from StarExec. An additional information should be shown when such data is about to be fetched. The desired effect is that each request is responded as early as possible. Also every additional calculation which is repeatedly needed should be cached, too.

6

Star-Exec-Presenter

The Star-Exec-Presenter is actually the concrete work of this thesis. As a REST-based web application it is an interface to the StarExec service and it provides a caching mechanism. This chapter will give detailed insights in the architectural approach to build the Star-Exec-Presenter according to the requirements.

Application Structure

Star-Exec-Presenter is written entirely in Haskell by utilizing its module system and the Yesod Web Framework. Detailled information are listed in chapter 7 "Implementation". In this section I will discuss the particular pieces that form Star-Exec-Presenter into a full server-side web application. To better understand the application structure, it could easily be interpreted as a Model-View-Controller (MVC) pattern where the Model stands for the data model, the View is represented by the templates which form the HTML presentation, and the Controller in the form of the several request handlers.

13

Data Model

The data model can be devided into two groups, the first containing the data types used for the REST-API which I will discuss later, and the second being the persistent data types of the actual data like the job results.

The first group of data types is used as identifiers for the second group, so they are directly related to each other. For instance there is `JobID` type that relates to `Job` type that containes the infos about a job as well as to a list of `JobResult` containing the actual results of the particular job. All these types mostly are unified types for the different entities of StarExec as well as imported entities.

To complete the example from above the type `JobID` has three different appearences, first being `StarExecJobID` representing a job on StarExec, the second being `LriJobID` representing an imported job from the 2007 Termination Competition and the third being `UibkJobID` for the imported jobs from Innsbruck. Additionally, a `Job` also has three different appearances (`StarExecJob`, `LriJob`, `UibkJob`) as well as `JobResult` (`StarExecResult`, `LriResult`, `UibkResult`).

The reason to have three different kinds of appearances for each type is because of having the data from StarExec in combination with the imported data from previous competitions. This old data differs in its form from that of the 2014 Termination Competition, so to consider these differences in a safe way I decided to isolate them persistent-wise. So, for each kind of data there are three database tables used to store that data.

Code-wise, the data model can be found within the `Presenter.Model` module, which combines all the groups of data types mentioned above as well as the persistent data types, which are defined in the `config/models` file.

Templates

The templates are essentially the view of Star-Exec-Presenter, they are filled with the requested information in a proper way to be displayed via the web browser. HTML[1] is very well suited for this purpose. Other types of data formats like JSON[2], CSV[3] or XML[4] can be considered as well but aren't a subject of matter, although they are used in the communication with the StarExec service.

---

[1] Hypertext Markup Language, a markup language designed to be interpreted by web browsers
[2] Javascript Object Notation, a compact subset of javascript designed for data exchange
[3] Comma-Separated Values, a data format to represent tables or list
[4] Extensible Markup Language, a data format for hierarchically structured data like trees

All templates can be found in the `templates` folder, some are located within the `Presenter.Utils` module.

Request Handlers

The controller part of Star-Exec-Presenter is represented by the several request handlers. Each handler responds to a certain URL[5] requested by a client application. For instance there are handlers that only return the information of a particular entity like a solver or a job, there are also handler that initiate specific requests to StarExec or display the results of a competition. The request handlers are generally managing the logic behind Star-Exec-Presenter.

All handlers are implemented within the `Handler` module.

Caching

Fast responses, especially for handlers that can return a huge bunch of data, is a main requirement for Star-Exec-Presenter, so there has to be an effective solution. We came up with the simple idea of caching all data in the database. For instance, when the user requests the results of one or more certain jobs, the application first looks in the database for the respective data and returns it. If there isn't found any related data it will be fetched from StarExec. But this fetching can affect a fast response, so it will be separated from the actual response by running in the background.

This is achieved by starting a new thread within the Star-Exec-Presenter application. The following figure shows the process from a request to a response:

This database based caching mechanism has its downside when it comes the the competition results, as they are calculated based on the results of the related jobs. And as the data of these jobs is stored in the database, there is no need for the results to be stored either. So the results have to be cached in another way which is within the application's memory. For this purpose Star-Exec-Presenter uses STM[6].

The reason to cache the competition results is because it takes too long to fetch the related job results from the database and do the calculations, as a usual competition

---

[5]Uniform Resource Locator

[6]Software Transactional Memory, a mechanism for concurrent programming with shared resources

has a huge amount of data to process. For instance the 2014 Termination Competition has three meta categories with a total of 19 jobs that produced an overall amount of 37.880 job-pairs. That doesn't seem to be that much but one requirement was to have short response times. So we decided to offload the computation to a separate worker thread. The results are saved within the application's memory and are accessed via STM. The calculating worker thread is the only instance which writes in this cache, all request handlers can only read from it.

STM helps in this case because it can prevent deadlocks from multiple simultaneous access to the shared memory making the whole application more reliable. To achieve this goal every access to the shared memory (our cache of competition results) is being managed by transactions, much like transactions of a database. If two or more transactions try to access a resource simultaneously, all but one will be aborted. A mechanism to retry a rolled back transaction ensures that they will be finished anyhow. To achieve this behavior each transaction must be very small. The actual terminology is *atomic*. It is wise to have very small transactions because STM is not *fair* as it favors small and fast transactions over large and slow ones. STM doesn't work in FIFO[7] order. (Marlow 2013)

Star-Exec-Presenter itself only does the reading and writing process within STM. The calculation of the competition results is done outside of STM. Only in this way we can ensure that the transactions are very small – or *atomic*.

REST interface

Star-Exec-Presenter is a server-side web application which receives requests and responses with an appropriate HTML presentation. These requests have to point to a valid resource, a URL. A list of important resource-URLs of the application including their accepted HTTP methods follows:

```
/jobs/#JobID                    GET
/solvers/#SolverID              GET
/benchmarks/#BenchmarkID        GET
/pairs/#JobPairID               GET
/post_procs/#PostProcID         GET
/results/#Query/*JobIds         GET
```

---

[7]First In First Out

```
/proofs/#Text                   GET
/registered                     GET
/competitions/#CompetitionInfoId   GET
/control                        GET POST
/import                         GET POST
```

Of course, an explanation of these routes, as they are called, is helpful. The first five routes are self-explaining. They take a specific identifier and return the infos of the related entity, wether it's a job, solver, benchmark, job-pair or post-processor. All these routes have a counterpart which lists all related entities. These routes are accessed without any parameter at the end.

The Results-Route takes a list of Job-Identifiers and returns all of their results. A special part of this route is its first dynamic path piece, the query parameter which defines a special case for this route. If there is given a specific query then the job results will be filtered according to it. More on this topic can be read in the section "Querying".

The Proof-Route is directly related the the Job-Pair-Route. If a job-pair contains a proof, this route is meant to display it. The Registered-Route lists all of the participant of the 2014 Termination Competition including links to their configuration on StarExec. The Competition-Route displays the results for a specific competition. This route has a counterpart, too, which lists all (publicly available) competitions.

The Control-Route is an interface to start a competition wether it's a full or a small one. Requested with the default HTTP method GET this route only displays its interface form. If this form is submitted the POST method is requested and a competition will be started. The Import-Route is used to import old data from previous competitions. It has a small form with a select field and a file input to upload a file. It takes a zip file with the content which has to be defined by the select field.

There are other routes that give further information as well as some routes that are older versions of the aforementioned ones. They are still accessible for legacy purposes and redirect to the new ones.

As Star-Exec-Presenter is meant to be a simple interface for StarExec and the Termination Competition in the first place, these routes are quite simple as well. They only accept GET or POST requests and try to work according to principles of REST. REST is an acronym for Representational State Transfer and was first mentioned by Roy Fielding in his dissertation "Architectural Styles and the Design of Network-based Software Architectures". REST respects the stateless nature of HTTP and

defines that every resource has its unique unified resource identifier (URI). (Fielding 2000)

HTTP methods like GET, POST, PUT or DELETE are used to get, create, change or remove the specific resource. The GET method is nullipotent, which means that a GET request never has any side effects and never changes anything respectively. POST, PUT and DELETE requests are idempotent, resulting in the same outcome no matter how much they are called. The URI, e.g. `/clear_the_database` itself never causes a change to the data unless it is requested with one of the mentioned HTTP methods. Apart from that, this is just an example which isn't recommendable for a real world REST API.

REST utilizes the available web technology for this purpose and isn't based on any standard like SOAP[8]. In contrast to SOAP REST is more like an architectural strategy or design pattern to develop a web application, whereas SOAP is a protocol. REST relies on the hypertext transfer protocol (HTTP) to mediate the request and simplifies the development process with it. To explain this principle a little further I want to emphasize the aforementioned MVC design pattern. In this context a REST interface can be considered the controller of the application which responses to the request with a suitable presentation.

A suitable presentation is HTML as well as XML or JSON. Star-Exec-Presenter currently only sends HTML. But this behavior could be altered by using a special header in the HTTP request. So, to get the data as XML or JSON, the Accept-Header could be set to the values `Accept: text/xml` or `Accept: text/json`. The web application then can evaluate this header to respond either with an XML or with a JSON presentation. (Richardson 2007)

It is widely accepted not to use verbs as parts of the URL rather than plural nouns. For instance the route `/get_job_results/#JobID` could be better written as `/job_results/#JobID` or simply `/results/#JobID`. Also, it is recommended to have a generalized route, e.g. `/jobs`, which lists all resources, whereas a parameter added to the URL results in returning only the specified resource. In the case of Star-Exec-Presenter especially the Results-Route is a special case takes an arbitrary amount of `JobIDs`.

Querying

---

[8]Simple Object Access Protocol, a web protocol for structured information

7

Implementation

The Star-Exec-Presenter is a web application written in Haskell and based upon the Yesod Web Framework. In this chapter I will talk about the actual implementation of the tool as well as about the benefits of using especially Haskell and Yesod.

Haskell

"Haskell is a general purpose, purely functional programming language incorporating many recent innovations in programming language design. Haskell provides higher-order functions, non-strict semantics, static polymorphic typing, user-defined algebraic datatypes, pattern-matching, list comprehensions, a module system, a monadic I/O system, and a rich set of primitive datatypes, including lists, arrays, arbitrary and fixed precision integers, and floating-point numbers. Haskell is both the culmination and solidification of many years of research on non-strict functional languages." (Peyton Jones 2003)

This quote stems from Simon Peyton Jones and his paper "Haskell 98 language and libraries: the revised report" and was used as well by Simon Marlow in his paper "Haskell 2010 Language Report". (Marlow 2010) It does sum up the important features of the Haskell programming language. In addition I want to add that Haskell has type inference and its type system is very expressive. In the following I try to explain the listed terms.

Haskell is a general purpose programming language which means it can be used to develop a web-server, a desktop application as well as a simple commandline tool. It is suited for a wide range of application domains. Haskell is a purely functional programming language, that is, it follows the functional programming paradigm and every pure function has no side effect. Functional programming relies on a mathematical approach where every calculation is defined by expressions. Everything is an expression, wether it's a function or a value. Purely means, that a function only works with its input arguments and returns the same result no matter how often it is called with the same arguments. Every expression is immutable so once they are defined they cannot be changed anymore. In contrast, imperative programming languages allow mutability and functions are subroutines that can have side effects. A side effect for instance can be a simple tracing that doesn't effect the function's result or it can be the change of a global state.

```haskell
-- example of a pure function
mul a b = a * b


-- example of an impure function
mulM a b = do
    putStrLn ("multiplying " ++ (show a) ++ " with " ++ (show b))
    return (a * b)
```

Haskell has higher-order function and emphasizes their usage. Higher-order function are functions that require functions as input parameters or return a function. Therefore functions are data that can be passed round. Thanks to currying all functions are higher-order functions and thereby expressions. Currying means that a function always returns a new function if it isn't called with the full list of arguments.

```haskell
-- example for a higher-order function
map :: (a -> b) -> [a] -> [b]
map _ []     = []                 -- special case for an empty list
map f (x:xs) = f x : map f xs     -- f applied to the list's head
                                  -- recursive call of map with
                                  -- the list's rest
```

The example from above also shows pattern-matching, a rich programming technique to identify certain cases for a given expression. The map function from above

has two implementations, one for the case of the empty list which also is the base case for the recursion, and the main implementation that declares the algorithm of the function. Pattern matching helps simplifying code by reducing the usage of if-statements and alike.

To stay with the `map` function, list comprehensions are another way to define this method as well as lists in general. They are similar to set comprehensions in mathematics. For example the list comprehension `[ 2*x | x <- [1..10] ]` is equivalent to the set comprehension $\{2x | x \in \mathbb{N}, x \leq 10\}$. (Lipovača 2012) Both expressions result in a list of even natural numbers from two to 20. Referecing the `map` function from above, an alternative implementation follows:

```
map f xs = [ f x | x <- xs ]
```

Another example of list comprehensions is the solution to the first problem of Project Euler (https://projecteuler.net/problem=1). The exercise is to find the sum of all multiples of three or five below 1000:

```
sum [x | x <- [3..999], x `mod` 3 == 0 || x `mod` 5 == 0]
```

Another essential feature of Haskell is the non-strict semantics. Non-strict in terms of Haskell means that every expression is evaluated by need. So, only declaring an expression doesn't invoke its evaluation until it's needed, e.g. for output. For instance the fibonacci sequence can be seen as an infinite list. This sequence can also be implemented via a list comprehension:

```
let fibs = 0 : 1 : [ a + b | (a, b) <- zip fibs (tail fibs)]
```

Here, we have the declaration of an infinite list, but which will be evaluated only up until its seventh position. These non-strict semantics in Haskell are implemented as *lazy evaluation* which leeds to the expression from above also being evaluated by need. So until it is really needed, `fibs` is only a thunked[1] expression. But this strategy can lead to high memory usage especially for complex algorithms, so in some cases it is advised to use strict evaluation which can be used with the `seq` function in Haskell's `Prelude` module. Other modules with strict evaluations for example are `Data.Map.Strict` or `Data.List` the latter having a strict implementation of the

---

[1] a value that is yet to be evaluated

`Prelude`'s `foldl` function. Especially this function can lead to a stack overflow with a large list.[2]

Haskell has a rich type system which can automatically infer the type of an expression. For example the fibonacci sequence from above has the type `fibs :: Num b => [b]` which was inferred bei the `(+)` operator's type (`(+) :: Num a => a -> a -> a`). Not all types can be inferred and the compiler outputs a compiler error, so a type signature should be added. In most cases these signatures are optional but also can help improving the readability of the code. These signatures doesn't contain any special type information because they are polymorphic, that is, they can be used with any type that meets the functions requirements, may it be `Int` or `Double`. Polymorphic types are denoted by small alphanumeric characters like `a`, `b` or `t0`. They can be restricted by a type-class which is a set of functions. The `(+)` operator is a function of the `Num` type-class which is a set of functions to operate on numerical values. Type-classes are similar to interfaces of imperative programming languages like Java.

With Haskell a developer can declare custom algebraic datatypes. What that means is that a type can be defined by the algebraic operations *sum* and *product*, where a *sum* is a group of alternative data constructors and *product* is the combination of them. Below are some examples:

```haskell
-- sum of polymorphic types
data Either a b = Left a | Right b


-- product of polymorphic types
data Pair a b = Pair a b
```

The `Either` datatype is a special type that is used for computations that can result in an error, where the `Left` constructor contains the error and the `Right` constructor holds the successful result. This leads to Haskell's expressiveness which is one of the benefits of its type system. The `Either` datatype explains that a function can return an error, so a developer has to consider the returned value. Java in contrast is more verbose by forcing to add a `throws` statement to the method definition. A developer has to use a `try ... catch` block to manage an error. Another example for the expressive type system of Haskell is the `Maybe` datatype whose signature follows:

```haskell
data Maybe a = Nothing | Just a
```

---

[2]https://www.haskell.org/haskellwiki/Foldr_Foldl_Foldl%27

This datatype is used for computations that can have no result or that require optional arguments. The `Nothing` constructor means that there is no value, whereas the `Just` constructor wraps a result or an optional argument. The `Maybe` is usefull because it prevents `undefined` or `null` values via the type system. In Java, for example, careless developed programs can pass the compiler but also result in a `NullPointerException` whose cause can be difficult to resolve. Of course, that doesn't mean that a Haskell program doesn't have errors, but the type system and the compiler helps to reduce them or event prevent them all.

(O'Sullivan, Goerzen, and Stewart 2010)

Yesod Web Framework

(Snoyman 2012)

Routes

Templates

Persistence

8
Evaluation

Future considerations

- caching could be automated with a permanent background-worker which traverses the database's content after a given time period
- ajax-based implementation of the web-interface
- generalized export of all data

9
Summary

...

Glossary

References

Fielding, Roy Thomas. 2000. *Architectural Styles and the Design of Network-Based Software Architectures*. http://www.ics.uci.edu/~fielding/pubs/dissertation/top.htm.

Lipova□ a, Miran. 2012. *Learn You a Haskell for Great Good!: a Beginner's Guide*. San Francisco, CA: No Starch Press.

Marché, Claude, and Hans Zantema. 2007. "The Termination Competition." In *In Proc. RTA '07, LNCS 4533*, 303–13.

Marlow, Simon. 2010. "Haskell 2010 Language Report." https://www.haskell.org/definition/haskell2010.pdf.

———. 2013. *Parallel and Concurrent Programming in Haskell*. First edition. Beijing: O'Reilly.

O'Sullivan, Bryan, John Goerzen, and Donald Bruce Stewart. 2010. *Real World Haskell*. Auflage: 1. Sebastopol, CA: O'Reilly.

Peyton Jones, Simon L., ed. 2003. *Haskell 98 Language and Libraries: the Revised Report*. Cambridge, U.K. ; New York: Cambridge University Press.

Richardson, Leonard. 2007. *RESTful Web Services*. Farnham: O'Reilly.

Snoyman, Michael. 2012. *Developing Web Applications with Haskell and Yesod*. 1st ed. Sebastopol, CA: O'Reilly.

StarExec. 2013. "About StarExec." https://www.starexec.org/starexec/public/about.jsp.

Waldmann, Johannes. 2014. "Termtools Termination Competition Data." http://lists.lri.fr/pipermail/termtools/2014-July/000979.html.

Zantema, Hans. 2000. *Termination of Term Rewriting*.